

Lecture 18

what we've done and what's to come

Announcements

- Please fill out course feedback if you haven't yet.
 - Canvas should show you a pop up, or
 - go to <http://evaluationkit.stanford.edu>, or
 - go to Axess -> Student -> Course -> Section Evaluations.
- Your feedback is very important.
- It will help make future iterations of the course better!

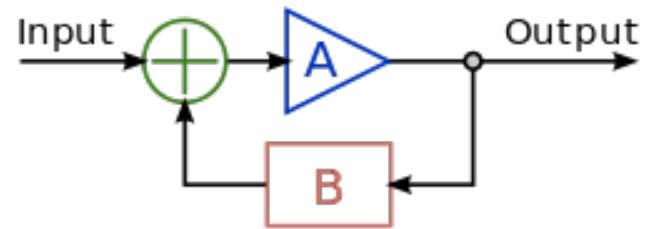


Figure 1: Feedback

Today

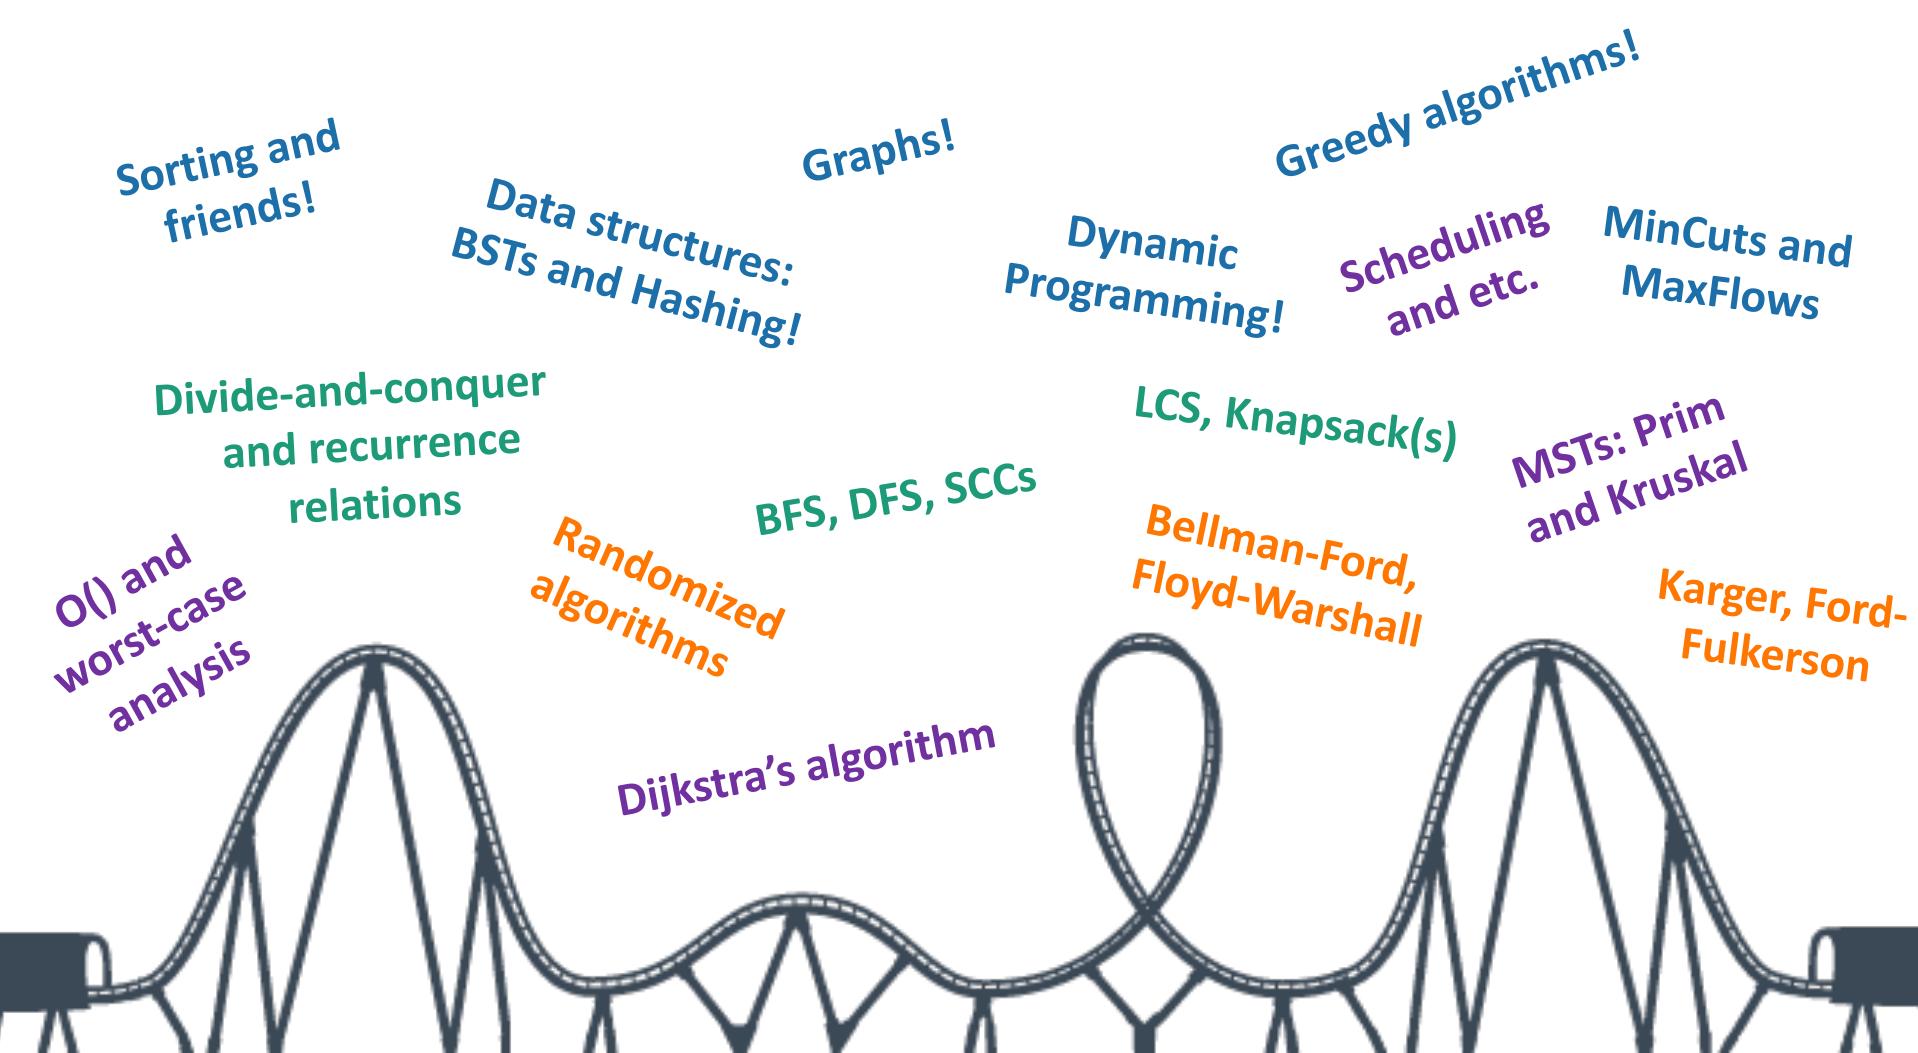
- What just happened?
 - A whirlwind tour of CS161



- What's next?
 - A few gems from future algorithms classes



It's been a fun ride...



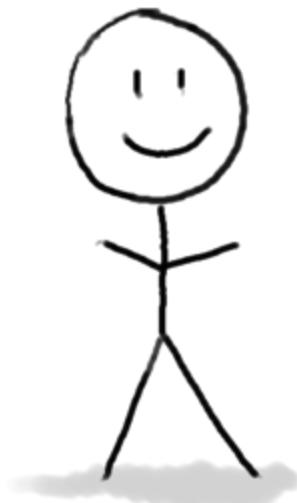
What have we learned?

17 lectures in 12 slides.

General approach to algorithm design and analysis

Can I do better?

To answer this question we need both **rigor** and **intuition**:



Algorithm designer



Plucky the
Pedantic Penguin

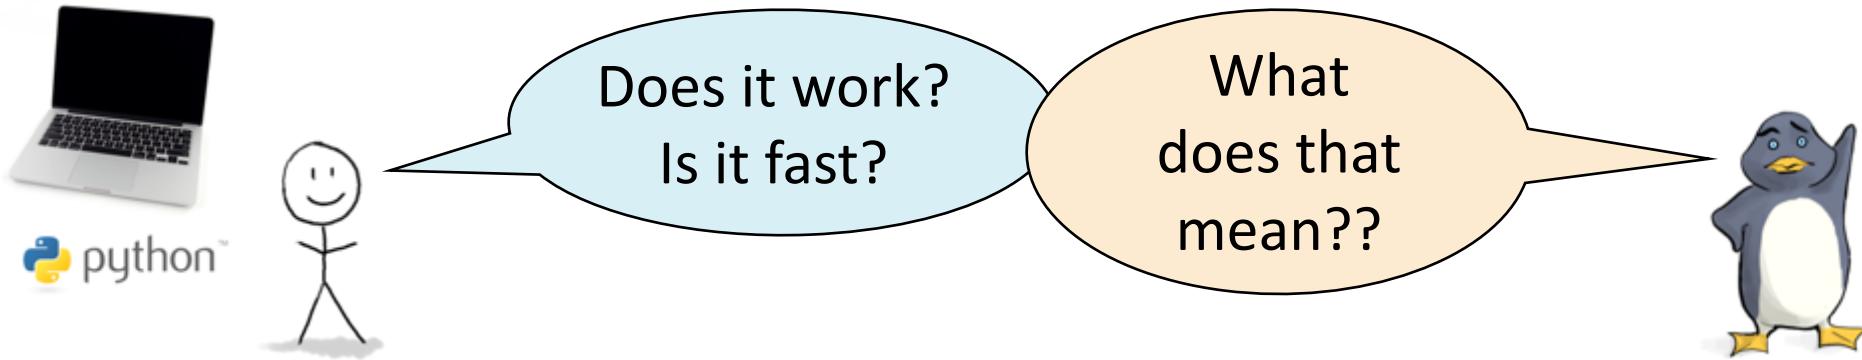
Detail-oriented
Precise
Rigorous



Lucky the
Lackadaisical Lemur

Big-picture
Intuitive
Hand-wavey

We needed more details

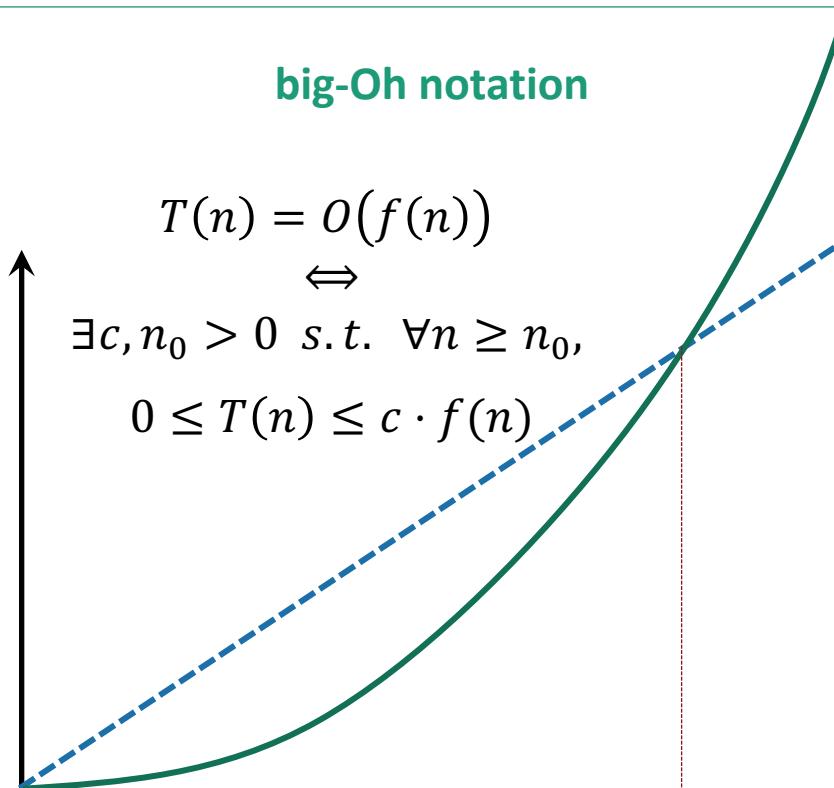


Worst-case analysis



big-Oh notation

$$T(n) = O(f(n)) \iff \exists c, n_0 > 0 \text{ s.t. } \forall n \geq n_0, 0 \leq T(n) \leq c \cdot f(n)$$



Algorithm design paradigm: divide and conquer

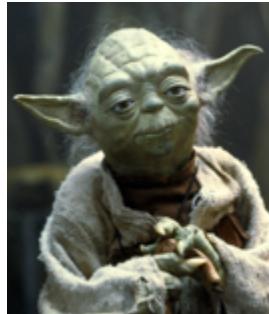
- Like MergeSort!
- Or Karatsuba's algorithm!
- Or SELECT!
- How do we analyze these?

By careful analysis!

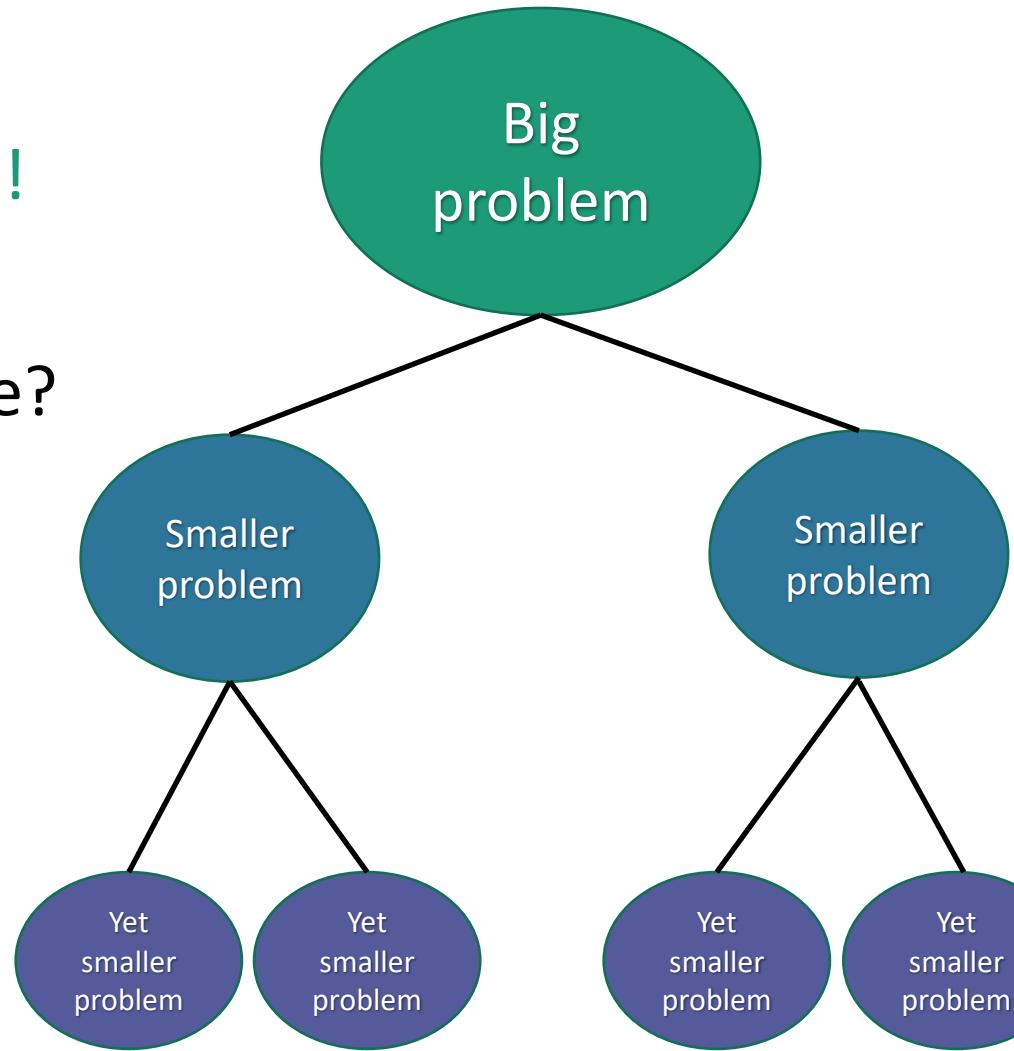


Plucky the Pedantic Penguin

Useful shortcut, the **master method** is.



Jedi master Yoda



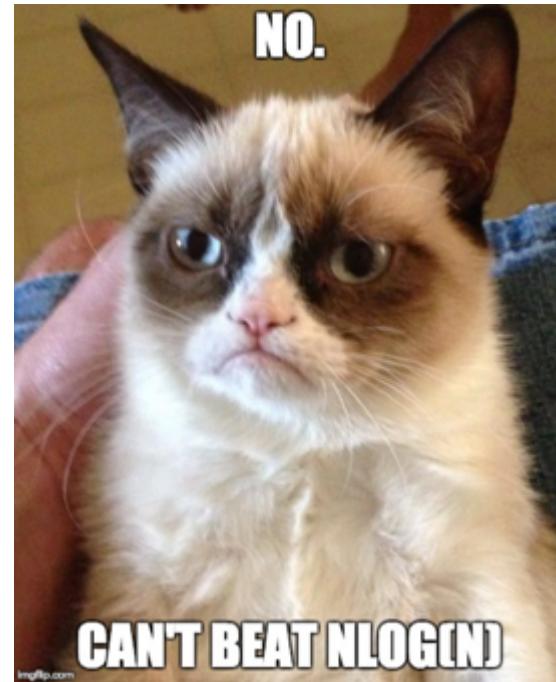
While we're on the topic of sorting Why not use randomness?

- We analyzed **QuickSort!**
- Still worst-case input, but we use randomness after the input is chosen.
- Always correct, usually fast.
 - This is a Las Vegas algorithm



All this sorting is making me wonder... Can we do better?

- Depends on who you ask:



- **RadixSort** takes time $O(n)$ if the objects are, for example, small integers!
- Can't do better in a **comparison-based** model.



beyond sorted arrays/linked lists: Binary Search Trees!

- Useful data structure!
- Especially the self-balancing ones!

Red-Black tree!

Maintain balance by stipulating that
black nodes are balanced, and that
there aren't too many **red nodes**.

It's just good sense!

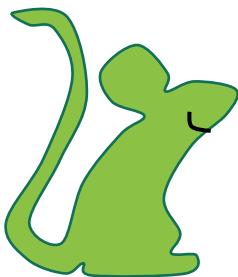


Another way to store things

Hash tables!

All of the hash functions
 $h:U \rightarrow \{1, \dots, n\}$

Choose h randomly from a universal hash family.

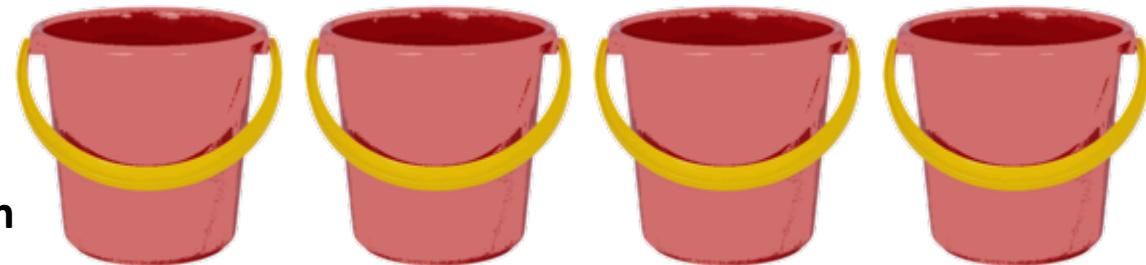


It's better if the hash family is small!
Then it takes less space to store h .



The universe

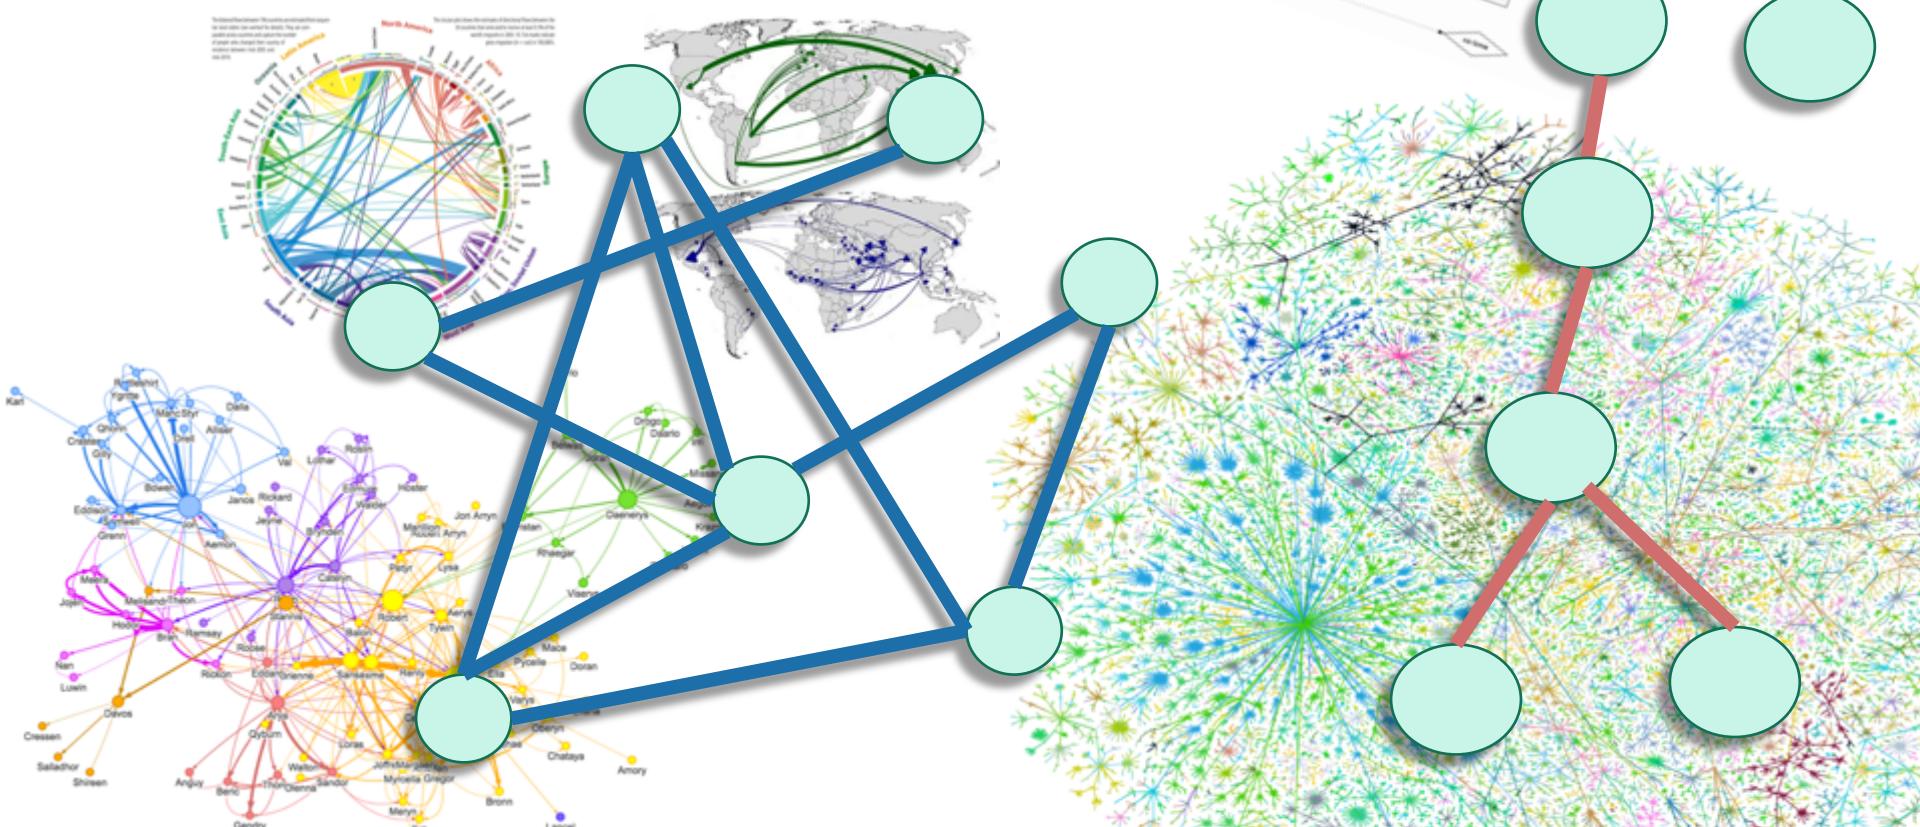
hash function h



Some buckets

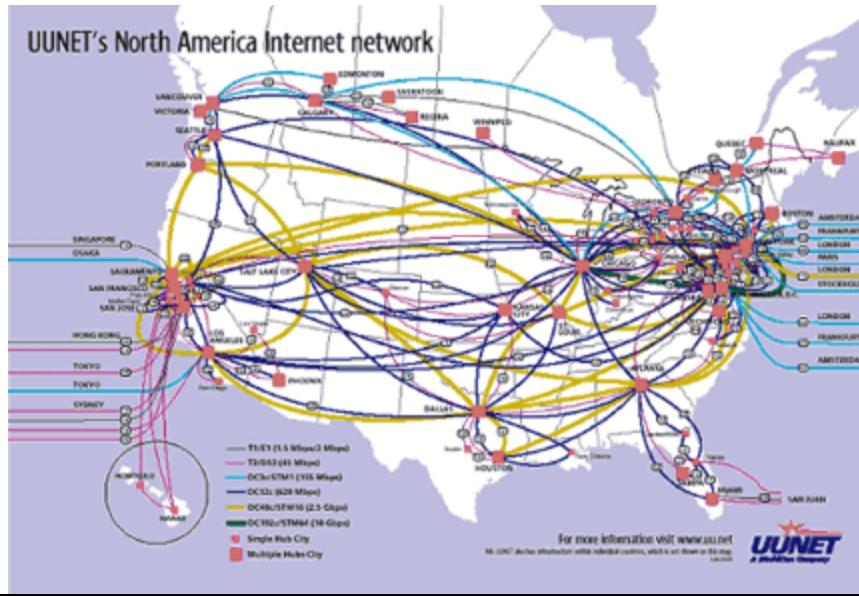
OMG GRAPHS

- BFS, DFS, and applications!
- SCCs, Topological sorting, ...



A fundamental graph problem: shortest paths

- E.g., transit planning, packet routing, ...
- Dijkstra!
- Bellman-Ford!
- Floyd-Warshall!



```
[DN0a22a0e3:~ mary]$ traceroute -a www.ethz.ch
traceroute to www.ethz.ch (129.132.19.216), 64 hops max, 52 byte packets
1 [AS0] 10.34.160.2 (10.34.160.2) 38.168 ms 31.272 ms 28.841 ms
2 [AS0] cwa-vrtr.sunet (10.21.196.28) 33.769 ms 28.245 ms 24.373 ms
3 [AS32] 171.66.2.229 (171.66.2.229) 24.468 ms 20.115 ms 23.223 ms
4 [AS32] hpr-svl-rtr-vlan8.sunet (171.64.255.235) 24.644 ms 24.962 ms 17.453 ms
5 [AS2152] hpr-svl-hpr2--stan-ge.cenic.net (137.164.27.161) 22.129 ms 4.902 ms 3.642 ms
6 [AS2152] hpr-lax-hpr3--svl-hpr3-100ge.cenic.net (137.164.25.73) 12.125 ms 43.361 ms 32.3 ms
7 [AS2152] hpr-i2-lax-hpr2-r&e.cenic.net (137.164.26.201) 40.174 ms 38.399 ms 34.499 ms
8 [AS0] et-4-0-0.4079.sdn-sw.lasv.net.internet2.edu (162.252.70.28) 46.573 ms 23.926 ms 17.453 ms
9 [AS0] et-5-1-0.4079.rtsw.salt.net.internet2.edu (162.252.70.31) 30.424 ms 25.770 ms 23.1 ms
10 [AS0] et-4-0-0.4079.sdn-sw.denv.net.internet2.edu (162.252.70.8) 47.454 ms 57.273 ms 73.1 ms
11 [AS0] et-4-1-0.4079.rtsw.kans.net.internet2.edu (162.252.70.11) 70.825 ms 67.809 ms 62.1 ms
12 [AS0] et-4-1-0.4070.rtsw.chic.net.internet2.edu (198.71.47.206) 77.937 ms 57.421 ms 63.6 ms
13 [AS0] et-0-1-0.4079.sdn-sw.ashb.net.internet2.edu (162.252.70.60) 77.682 ms 71.993 ms 73.1 ms
14 [AS0] et-4-1-0.4079.rtsw.wash.net.internet2.edu (162.252.70.65) 71.565 ms 74.988 ms 71.0 ms
15 [AS21320] internet2-gw.mx1.lon.uk.geant.net (62.40.124.44) 154.926 ms 145.606 ms 145.872 ms
16 [AS21320] ae0.mx1.lon.uk.geant.net (62.40.98.79) 146.565 ms 146.604 ms 146.801 ms
17 [AS21320] ae0.mx1.par.fr.geant.net (62.40.98.77) 153.289 ms 184.995 ms 152.682 ms
18 [AS21320] ae2.mx1.gen.ch.geant.net (62.40.98.153) 160.283 ms 160.104 ms 164.147 ms
19 [AS21320] swice1-100ge-0-3-0-1.switch.ch (62.40.124.22) 162.068 ms 160.595 ms 163.095 ms
20 [AS559] swizh1-100ge-0-1-0-1.switch.ch (130.59.36.94) 165.824 ms 164.216 ms 163.983 ms
21 [AS559] swiez3-100ge-0-1-0-4.switch.ch (130.59.38.109) 164.269 ms 164.370 ms 163.929 ms
22 [AS559] rou-gw-lee-tengig-to-switch.ethz.ch (192.33.92.1) 164.082 ms 170.645 ms 165.372 ms
23 [AS559] rou-fw-rz-rz-gw.ethz.ch (192.33.92.169) 164.773 ms 165.193 ms 172.158 ms
```



Bellman-Ford and Floyd-Warshall
were examples of...

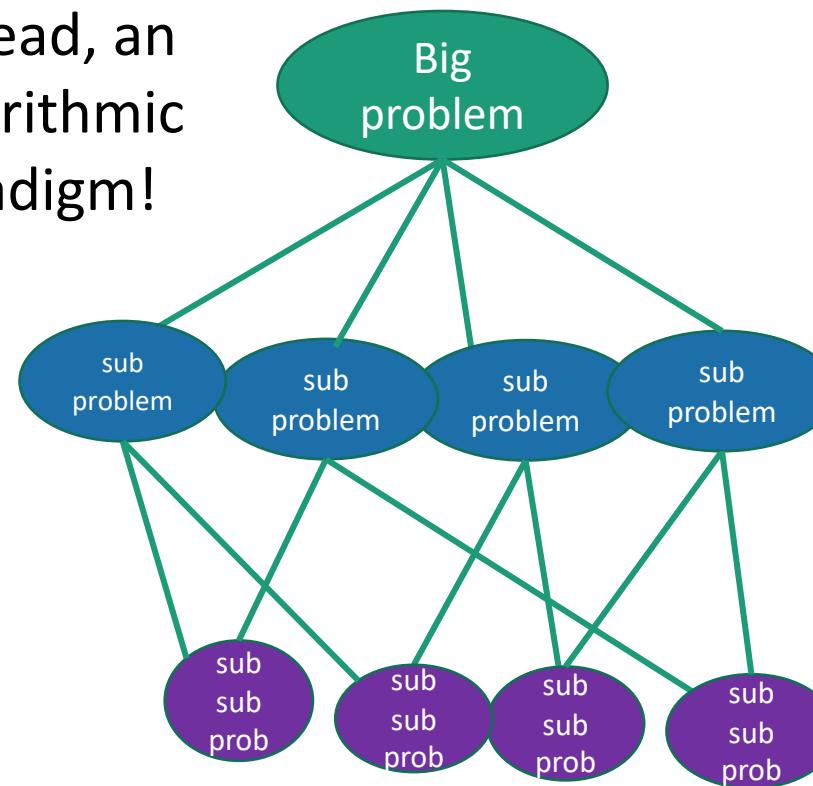
Dynamic Programming!

- Not programming in an action movie.



We saw many other examples, including Longest Common Subsequence and Knapsack Problems.

Instead, an algorithmic paradigm!



- **Step 1:** Identify optimal substructure.
- **Step 2:** Find a recursive formulation for the value of the optimal solution.
- **Steps 3-5:** Use dynamic programming: fill in a table to find the answer!

Sometimes we can take even better advantage of optimal substructure...with

Greedy algorithms

- Make a series of choices, and commit!



- Intuitively we want to show that our greedy choices never rule out success.
- Rigorously, we usually analyzed these by induction.
- Examples!

- Activity Selection
- Job Scheduling
- Huffman Coding
- Minimum Spanning Trees

Prim's algorithm:
greedily grow a tree

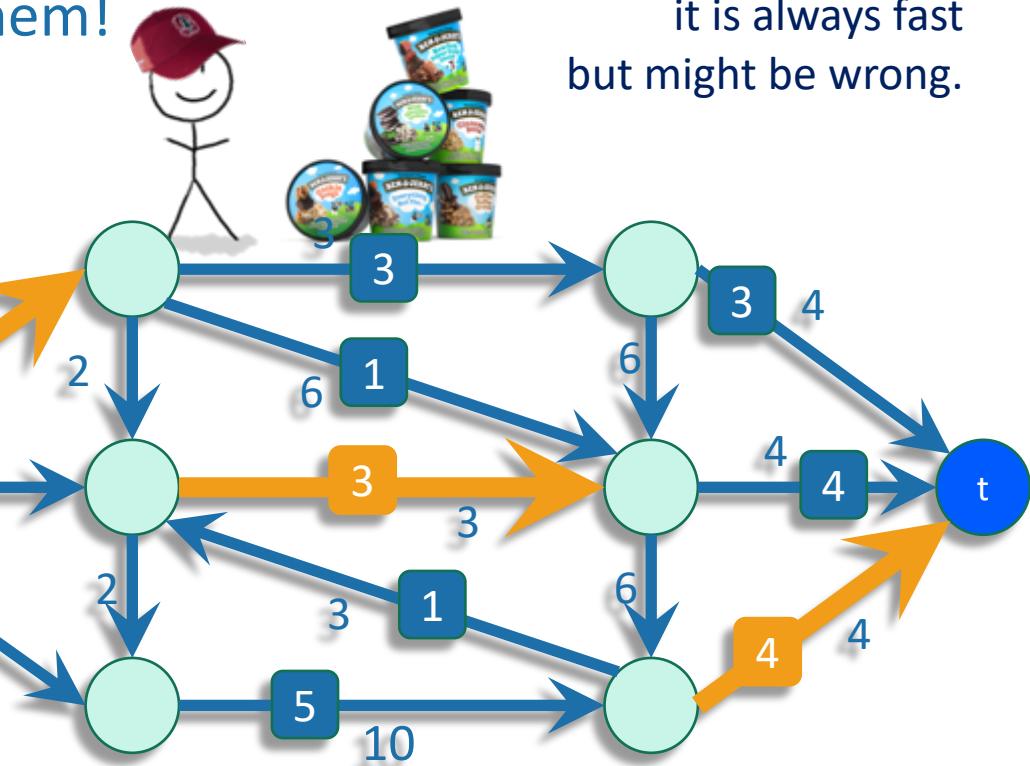
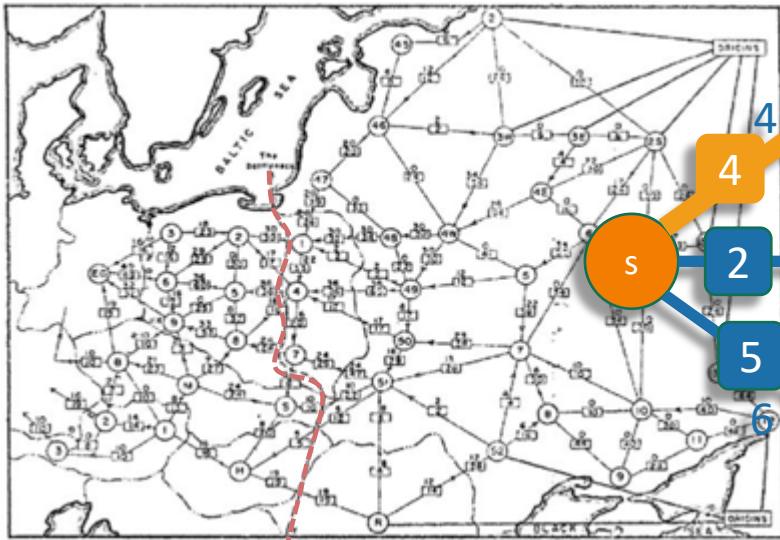


Kruskal's algorithm:
greedily grow a forest



Cuts and flows

- Global minimum cut:
 - Karger's algorithm!
- Minimum s-t cut:
 - is the same as maximum s-t flow!
 - Ford-Fulkerson can find them!
 - useful for routing
 - also assignment problems



Karger's algorithm is a Monte-Carlo algorithm:
it is always fast
but might be wrong.



And now we're here



What have we learned?

- A few algorithm design paradigms:
 - Divide and conquer, dynamic programming, greedy
- A few analysis tools:
 - Worst-case analysis, asymptotic analysis, recurrence relations, probability tricks, proofs by induction
- A few common objects:
 - Graphs, arrays, trees, hash functions
- A LOT of examples!



What have we learned? We've filled out a toolbox

- Tons of examples give us intuition about what algorithmic techniques might work when.
- The technical skills make sure our intuition works out.



But there's lots more out there



- What's next???

L₁ E₁ T₁ S₁ P₃ L₁ A₁ Y₄



A taste of what's to come

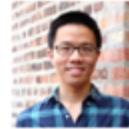
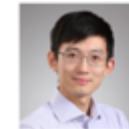
- CS154 – Introduction to Automata and Complexity
- CS163 – The Practice of Theory Research
- CS166 – Data Structures
- CS168 – The Modern Algorithmic Toolbox
- MS&E 212 – Combinatorial Optimization
- CS250 – Error Correcting Codes
- CS252 – Analysis of Boolean Functions
- CS254 – Computational Complexity
- CS255 – Introduction to Cryptography
- CS259Q – Quantum Computing
- CS260 – Geometry of Polynomials in Algorithm Design
- CS261 – Optimization and Algorithmic Paradigms
- CS263 – Counting and Sampling
- CS265 – Randomized Algorithms
- CS269O – Introduction to Optimization Theory
- MS&E 316 – Discrete Mathematics and Algorithms
- CS352 – Pseudorandomness
- CS366 – Computational Social Choice
- CS368 – Algorithmic Techniques for Big Data
- EE364A/B – Convex Optimization I and II

findSomeTheoryCourses():

- go to theory.stanford.edu
- Click on “People”
- Look at what we’re teaching!

STANFORD THEORY GROUP

Faculty

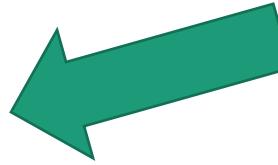
 Nima Anari	 John C. Mitchell	 Li-Yang Tan
 Dan Boneh	 Omer Reingold	 Gregory Valiant
 Adam Bouland	 Aviad Rubinstein	 Jan Vondrák
 Moses Charikar	 Amin Saberi	 Mary Wootters
 Ashish Goel	 Tsilei Schramm	
 Tengyu Ma	 Aaron Sidford	

...and many many more!

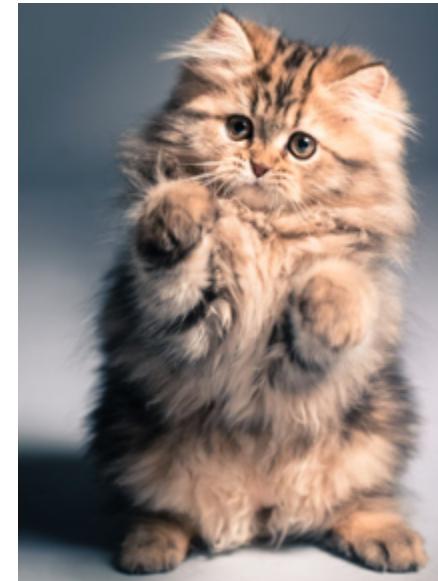
Today

A few gems

- Linear programming
- Random projections
- Low-degree polynomials



This will be pretty fluffy,
without much detail –
take more CS theory
classes for more detail!



Linear Programming

- This is a fancy name for optimizing a linear function subject to linear constraints.
- For example:

Maximize

$$x + y$$

subject to

$$x \geq 0$$

$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

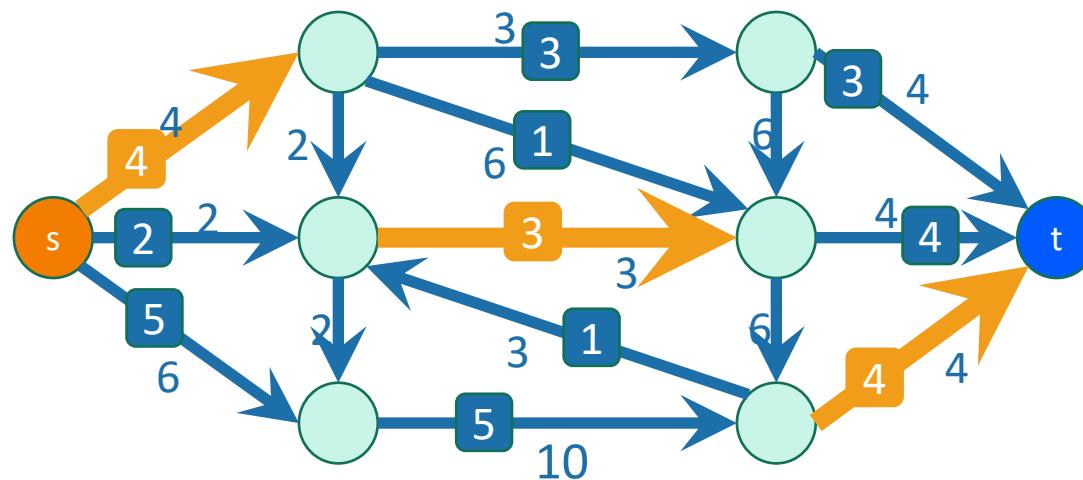
- It turns out the be an extremely general problem.

We've already seen an example!

Maximize
the sum of the
flows leaving s

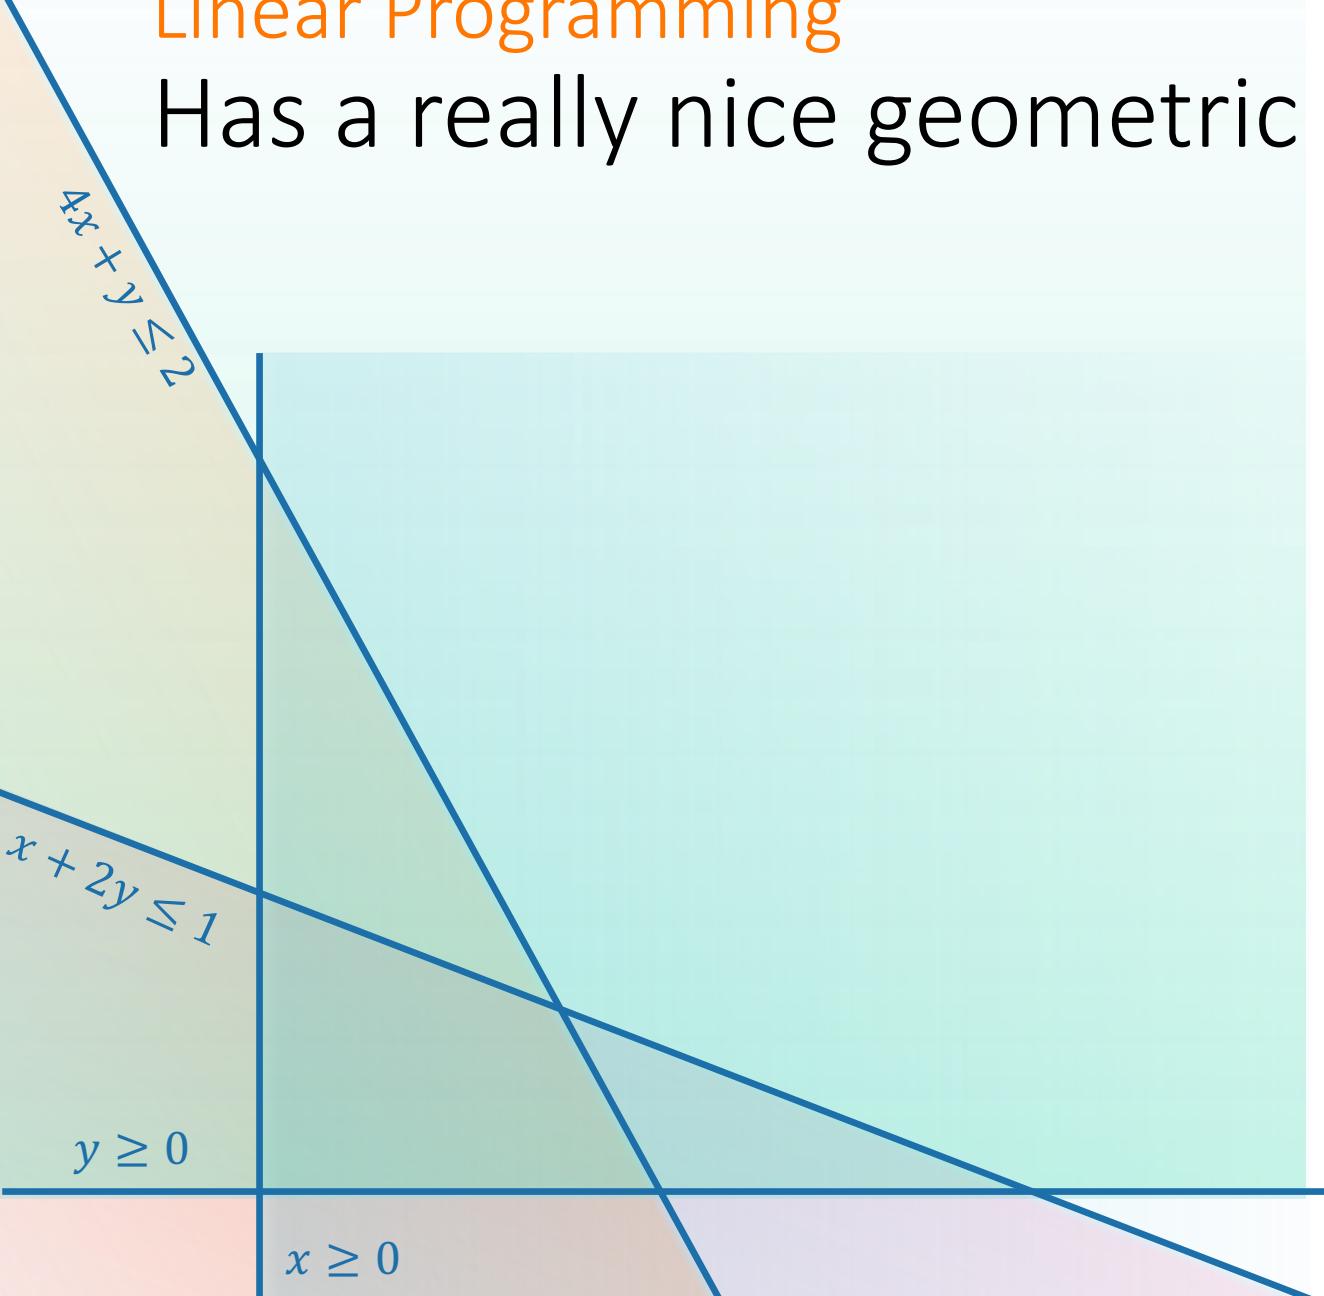
subject to

- None of the flows are bigger than the edge capacities
- At every vertex, stuff going in = stuff going out.



Linear Programming

Has a really nice geometric intuition



Maximize

$$x + y$$

subject to

$$x \geq 0$$

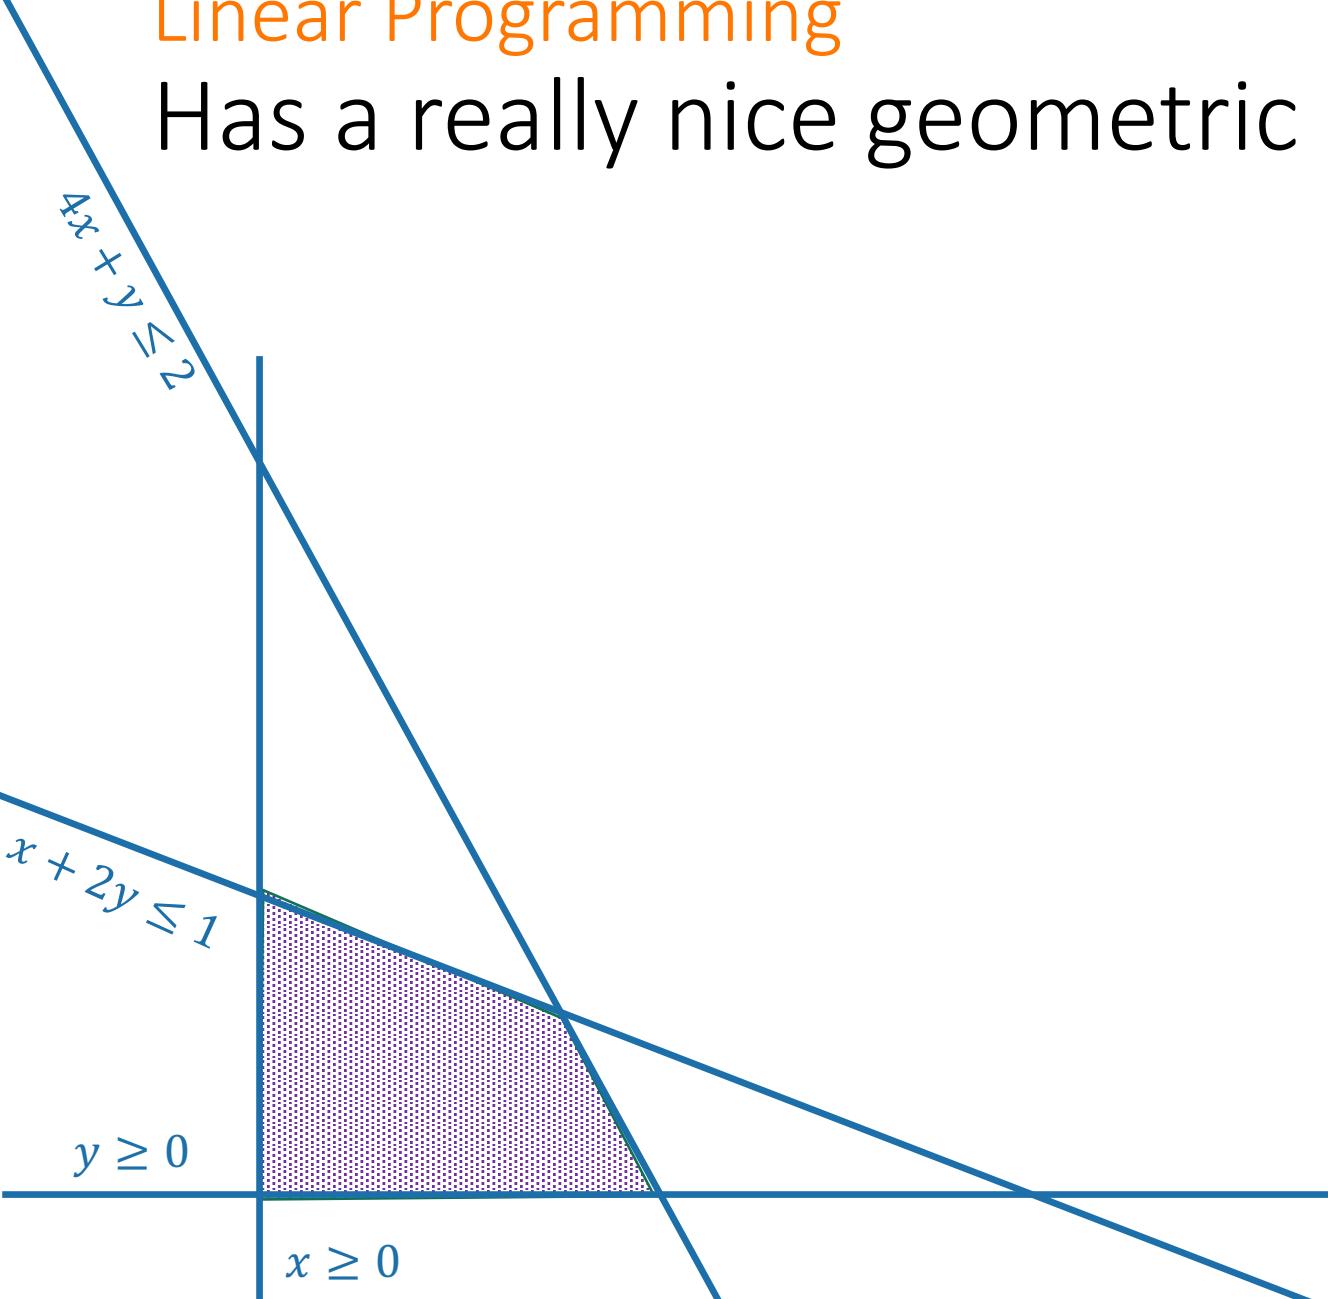
$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

Linear Programming

Has a really nice geometric intuition



Maximize

$$x + y$$

subject to

$$x \geq 0$$

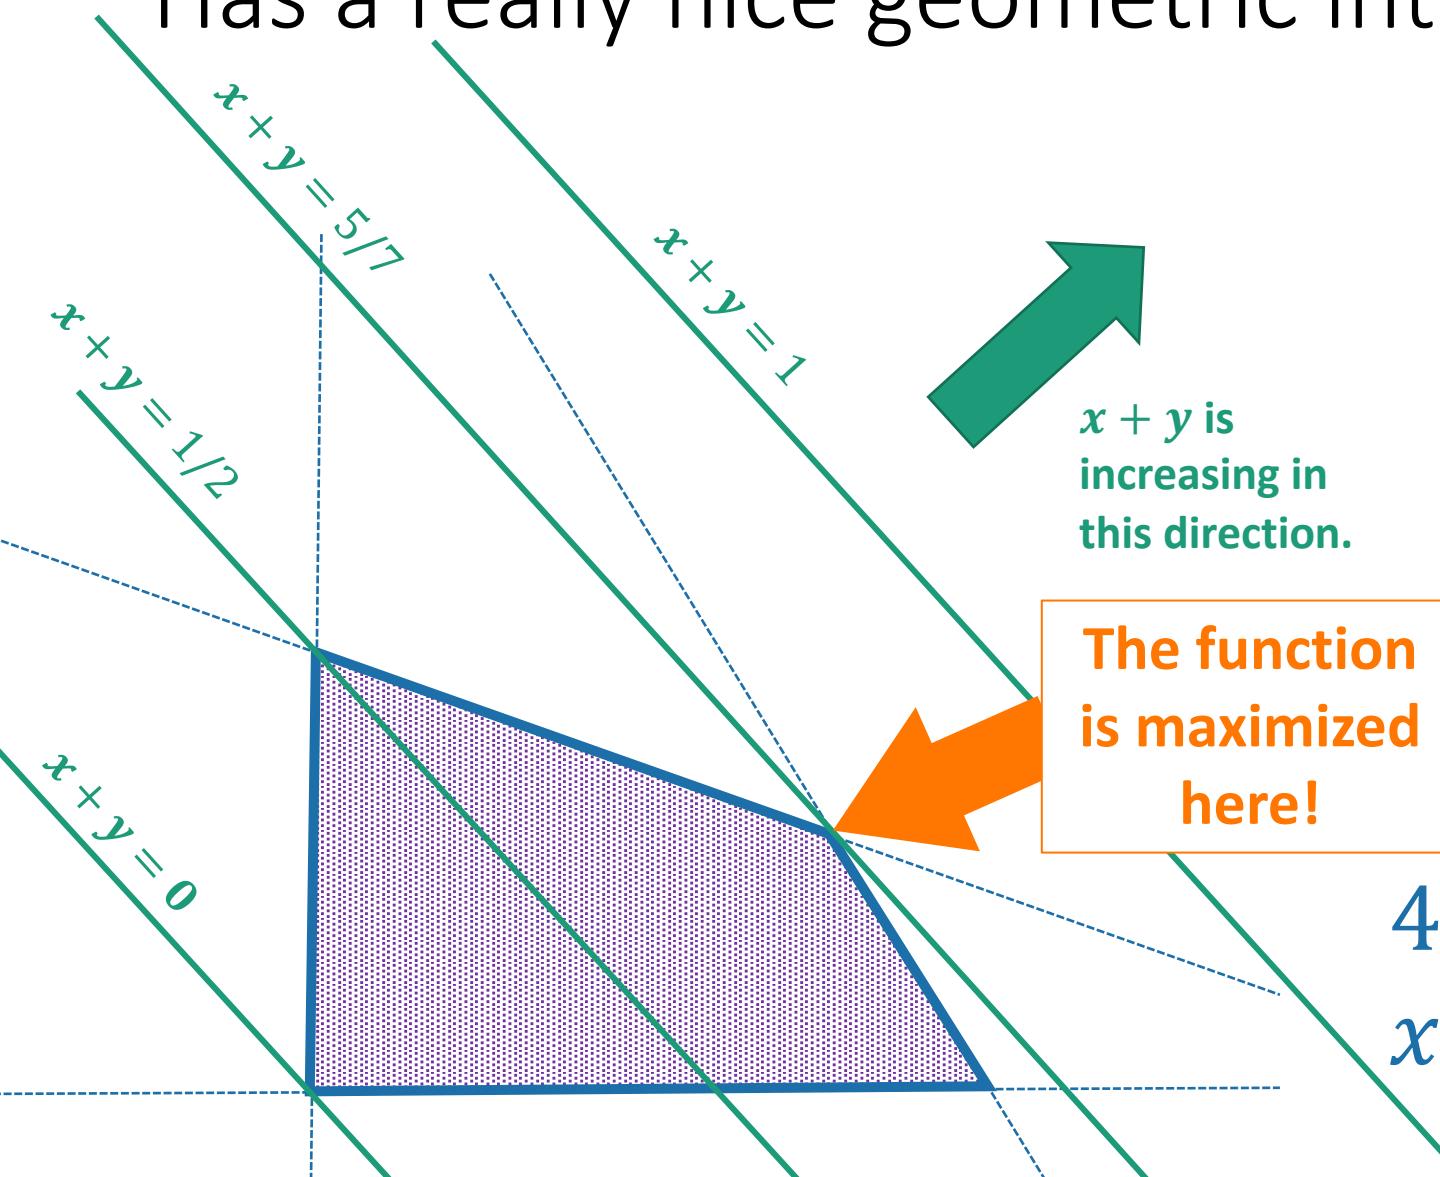
$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

Linear Programming

Has a really nice geometric intuition



Maximize

$$x + y$$

subject to

$$x \geq 0$$

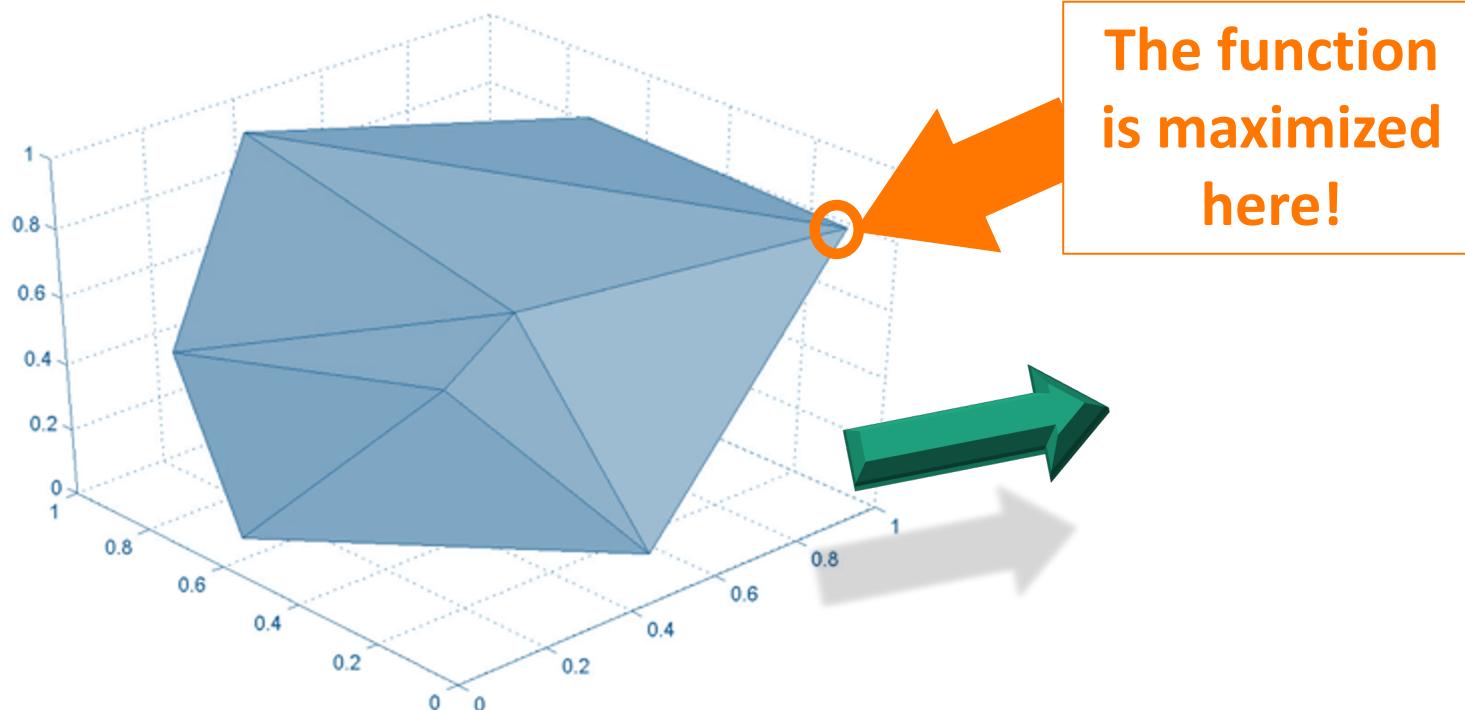
$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

In general

- The constraints define a **polytope**
- The function defines a **direction**
- We just want to find the vertex that is **furthest** in that direction.



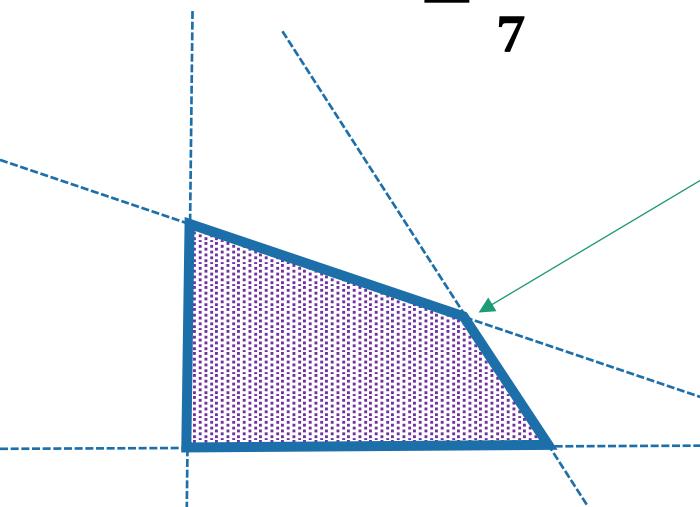
Duality

How do we know we have an optimal solution?

I claim that the optimum is $5/7$.

Proof: say x and y satisfy the constraints.

- $x + y = \frac{1}{7}(4x + y) + \frac{3}{7}(x + 2y)$
- $\leq \frac{1}{7} \cdot 2 + \frac{3}{7} \cdot 1$
- $= \frac{5}{7}$



You can check this point has value $5/7$...but how would we prove it's optimal other than by eyeballing it?

Maximize
 $x + y$

subject to

$$x \geq 0$$

$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

cute, but

How did you come up with $1/7, 3/7$?

I claim that the optimum is $5/7$.

Proof: say x and y satisfy the constraints.

- $x + y \leq (4x + y) + (x + 2y)$
- $\leq \cdot 2 + 1$
- $=$
- I want to choose things to put here
- So that I minimize this
- Subject to these things

Maximize

$$x + y$$

subject to

$$x \geq 0$$

$$y \geq 0$$

$$4x + y \leq 2$$

$$x + 2y \leq 1$$

Note: it's not immediately obvious how to turn that into a linear program, this is just meant to convince you that it's plausible.

In this case the dual is:

$$\begin{aligned} \min & 2w + z \text{ s.t. } w, z \geq 0, \\ & 4w + z \geq 1 \text{ and } w + 2z \geq 1 \end{aligned}$$

That's a linear program!

- How did I find those special values $1/7, 3/7$?
- I solved some linear program. Minimize the upper bound you get, subject to the proof working.
- It's called the **dual program**.

Maximize stuff
subject to stuff

Primal

The optimal values are
the same!

Minimize other stuff
subject to other stuff

Dual

We've actually already seen this too

The Min-Cut Max-Flow Theorem!

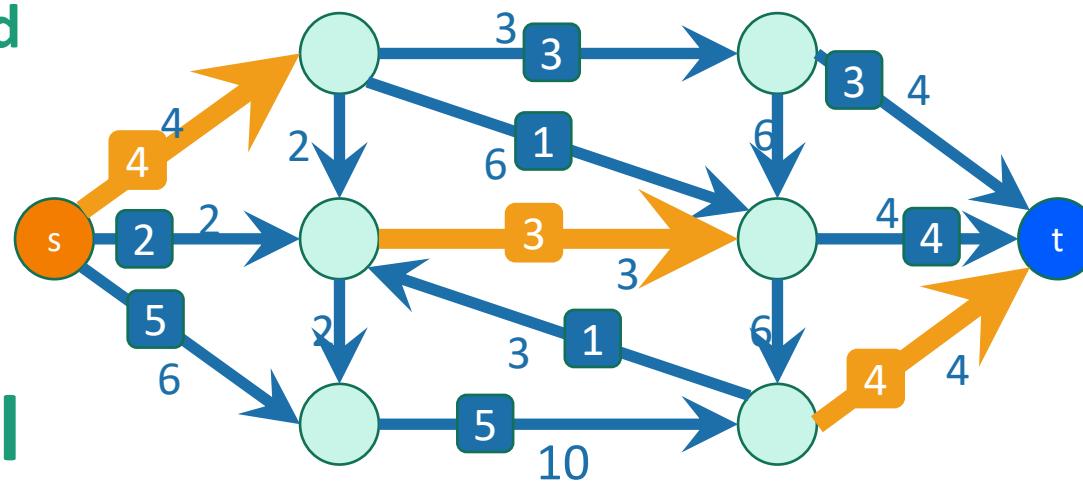
Maximize the sum of the flows leaving s
s.t.
All the flow constraints are satisfied

The optimal values are the same!

Minimize the sum of the capacities on a cut
s.t.
it's a legit cut

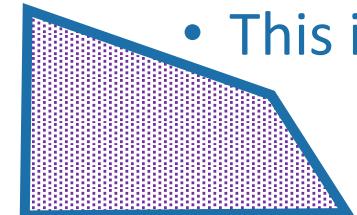
Primal

Dual



LPs and Duality are really powerful

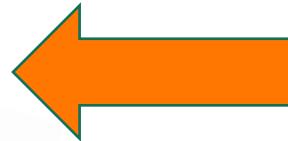
- This **general phenomenon** shows up all over the place
 - Min-Cut Max-Flow is a special case.
- Duality helps us reason about an optimization problem
 - The dual provides a **certificate** that we've solved the primal.
 - E.g., if you have a cut and a flow with the same value, you must have found a max flow and a min cut.
- We can solve LPs quickly!
 - For example, by intelligently bouncing around the vertices of the feasible region.
 - This is an **extremely powerful algorithmic primitive**.



Today

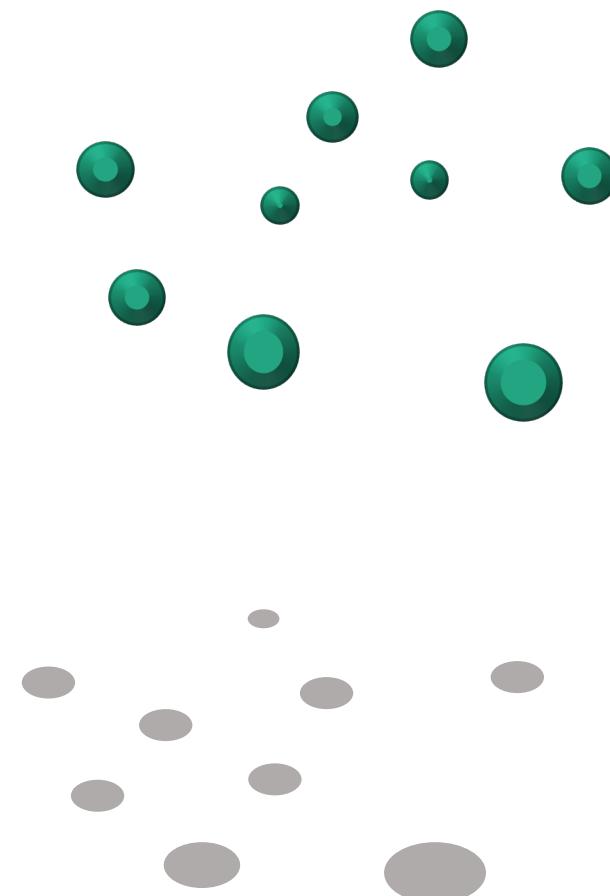
A few gems

- Linear programming
- Random projections
- Low-degree polynomials



A very useful trick

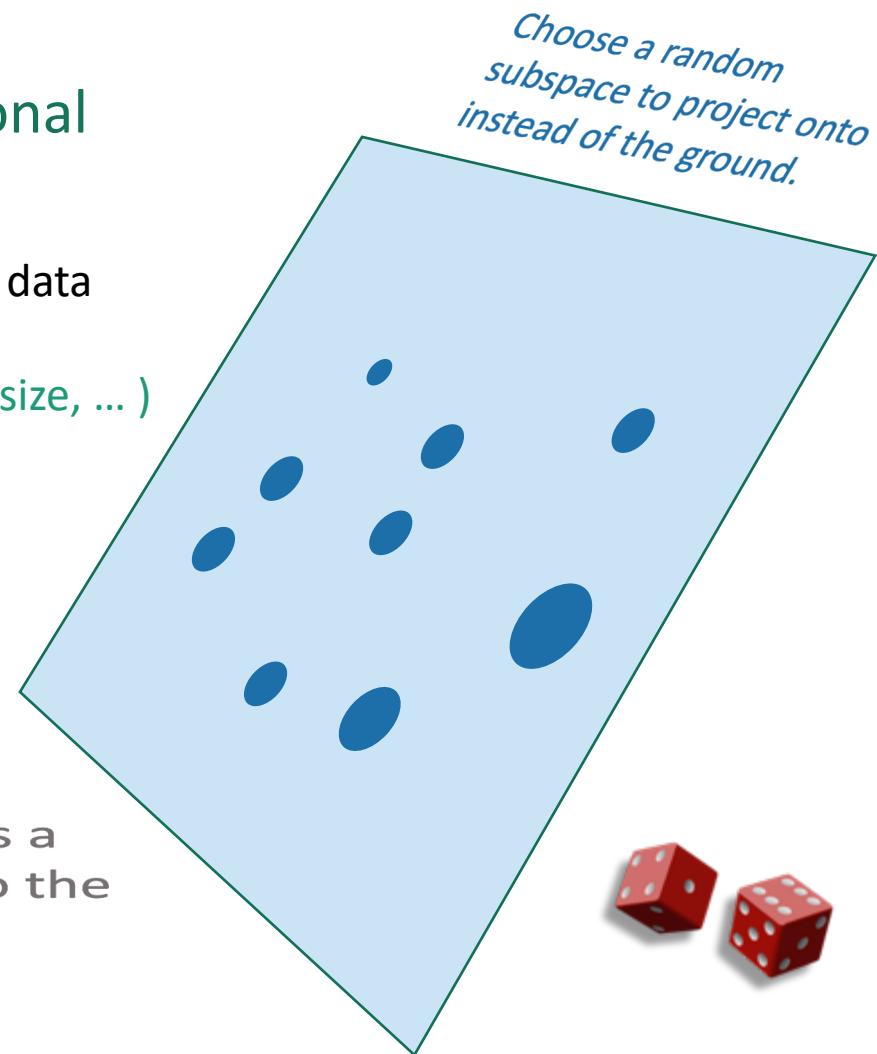
Take a random projection and hope for the best.



High-dimensional
set of points

For example, each data
point is a vector
(age, height, shoe size, ...)

Their shadow is a
projection onto the
ground.

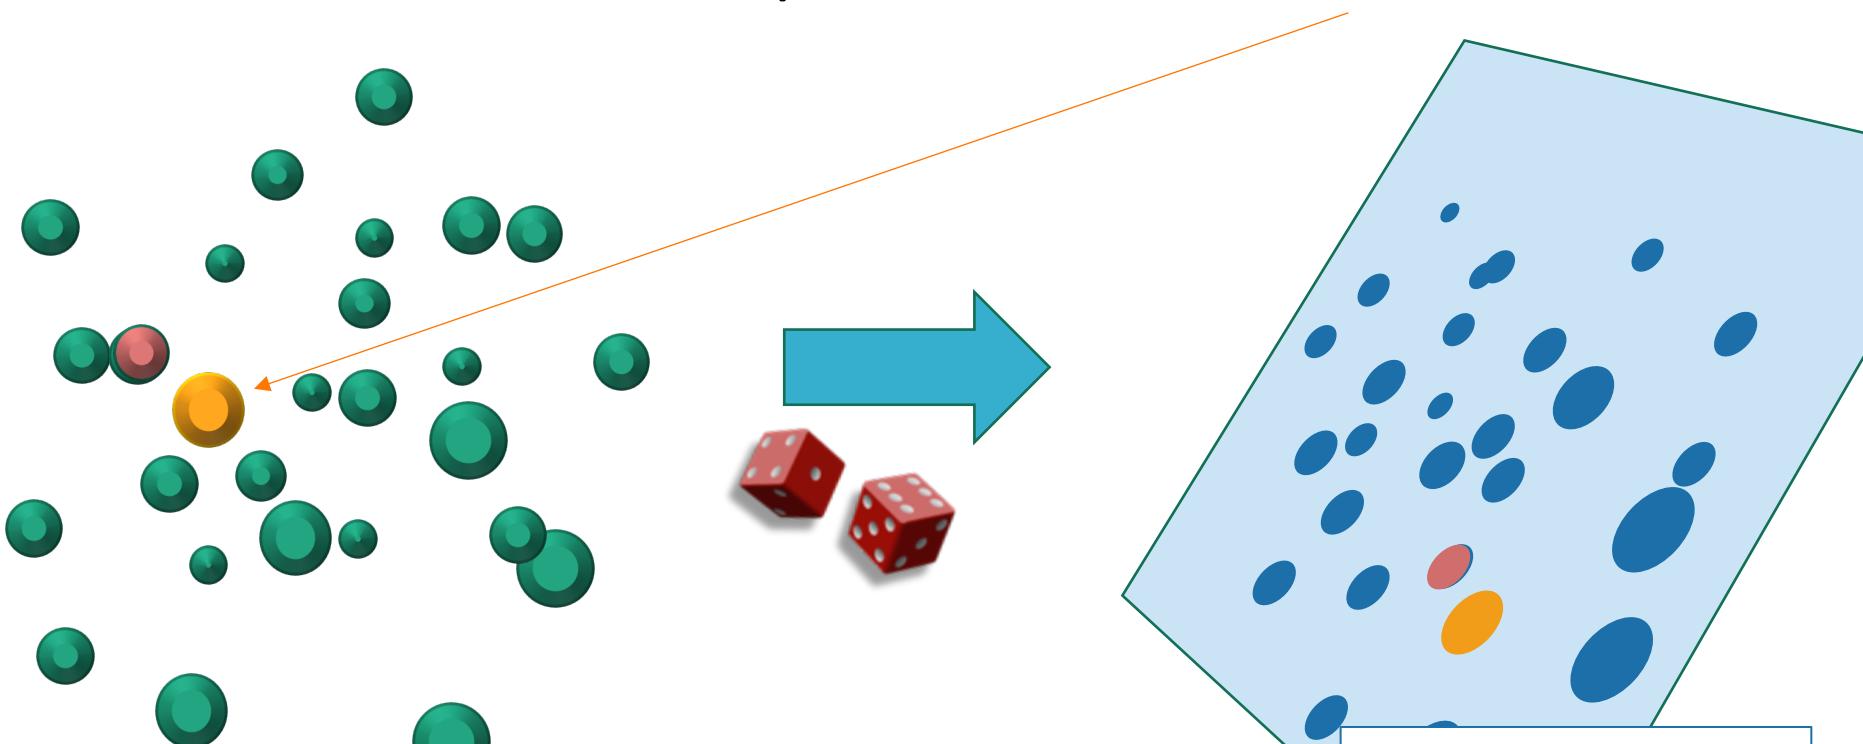


Why would we do this?

- High dimensional data takes a long time to process.
- Low dimensional data can be processed quickly.
- **“THEOREM”:** Random projections approximately preserve properties of data that you care about.

Example: nearest neighbors

- I want to find which point is closest to **this one**.



That takes a really long time in high dimensions.

Johnson-Lindenstrauss Lemma:
Euclidean distance is approximately preserved by random projections.

Find the closest point down here, you're probably pretty correct.

Another example: Compressed Sensing

- Start with a sparse vector
 - Mostly zero or close to zero

(5, 0, 0, 0, 0, 0.01, 0.01, 5.8, 32, 14, 0, 0, 0, 12, 0, 0, 5, 0, .03)

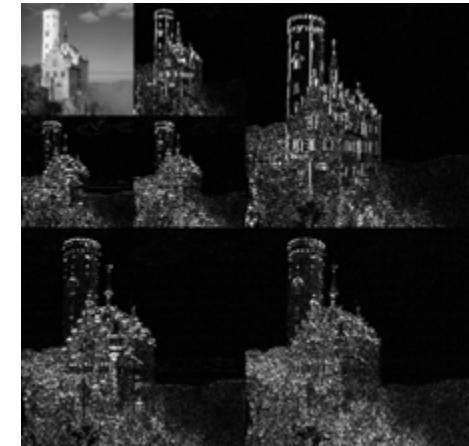
- For example:



This image is sparse

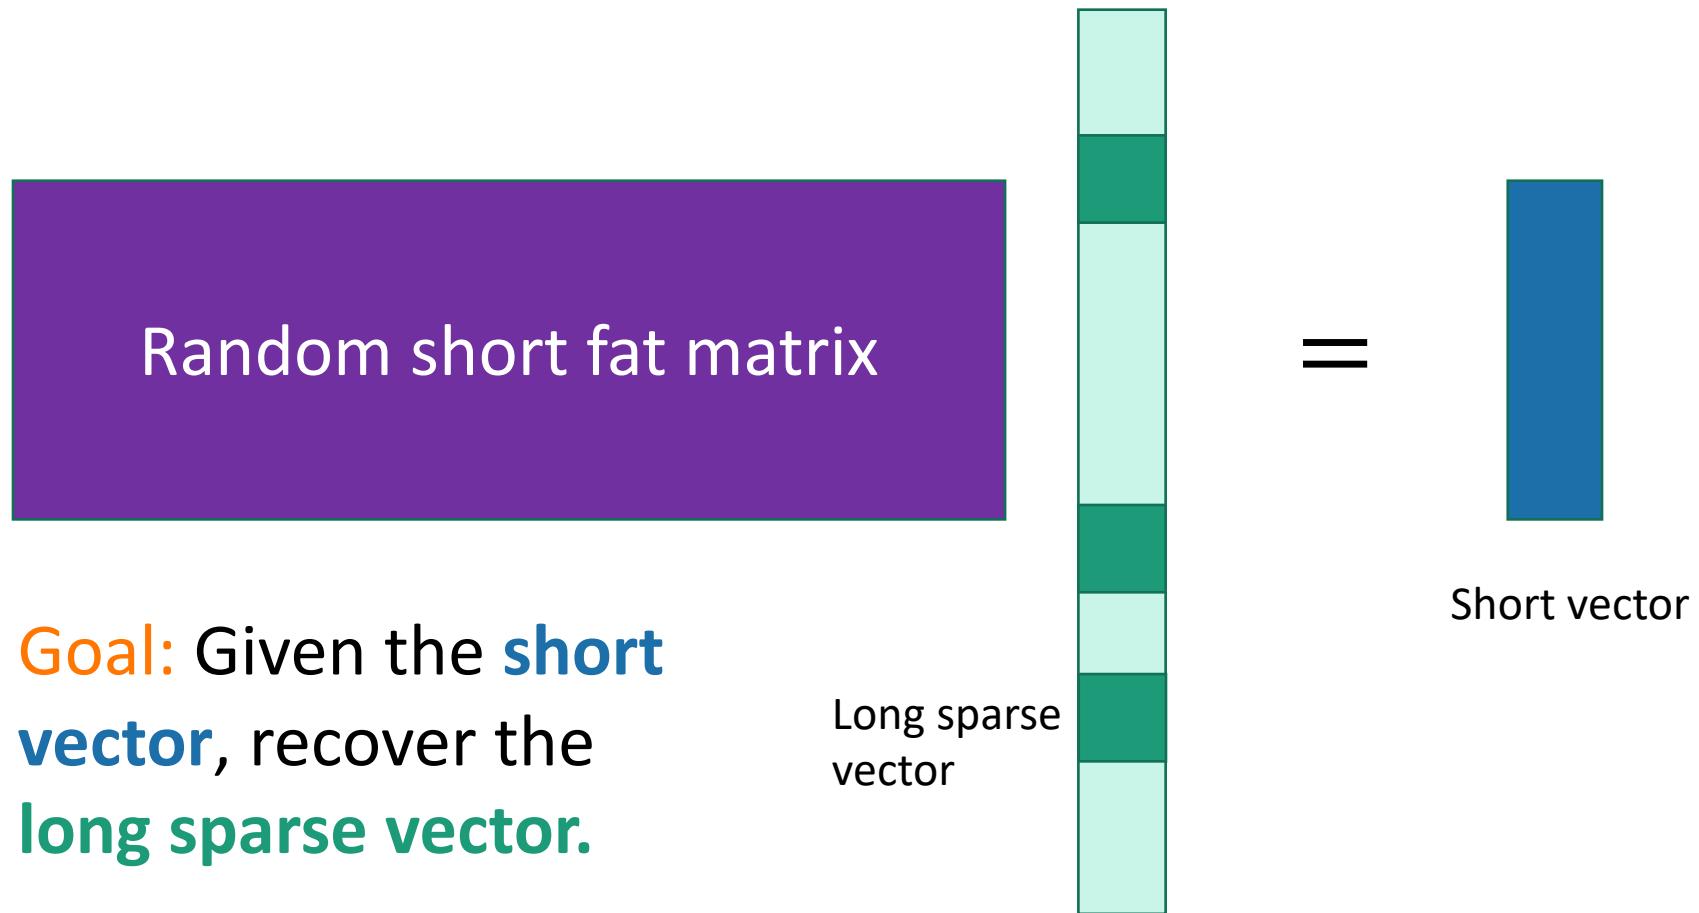


This image is sparse after I
take a wavelet transform.



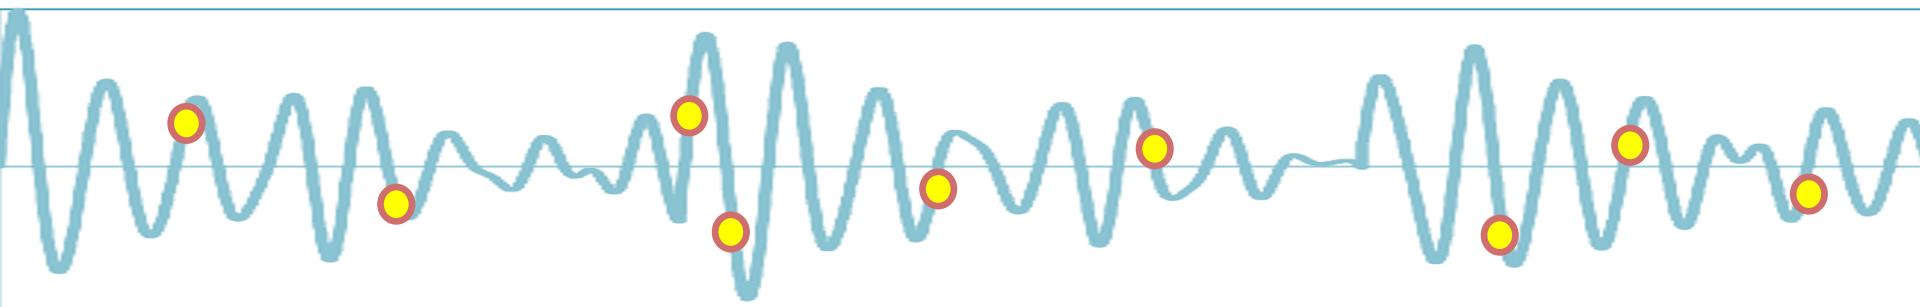
Compressed sensing continued

- Take a random projection of that sparse vector:

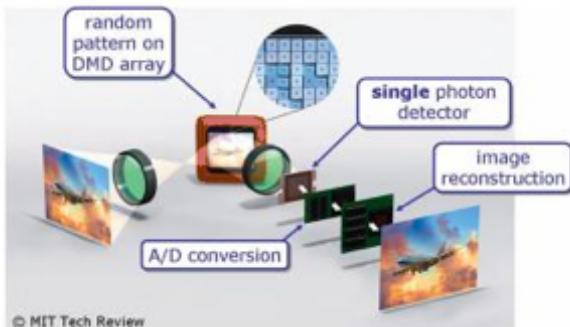


Why would I want to do that?

- Image compression and signal processing
- Especially when you **never have space to store the whole sparse vector to begin with.**



Randomly sampling (in the time domain) a signal that is sparse in the Fourier domain.

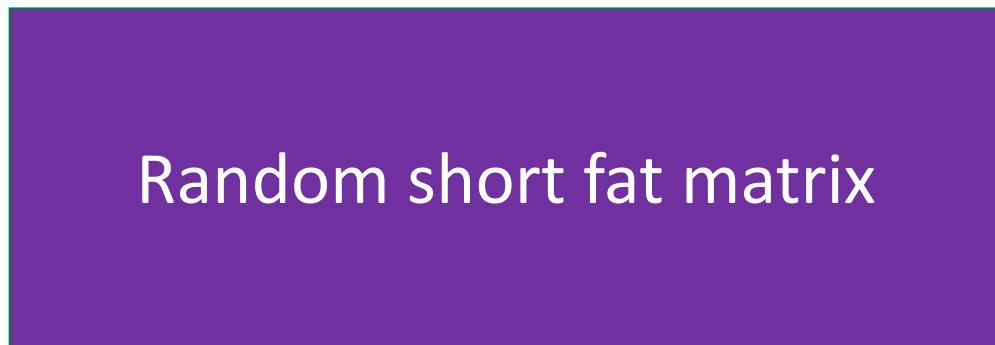


A “single pixel camera” is a thing.

Random measurements in an fMRI means you spend less time inside an fMRI



All examples of this:



=

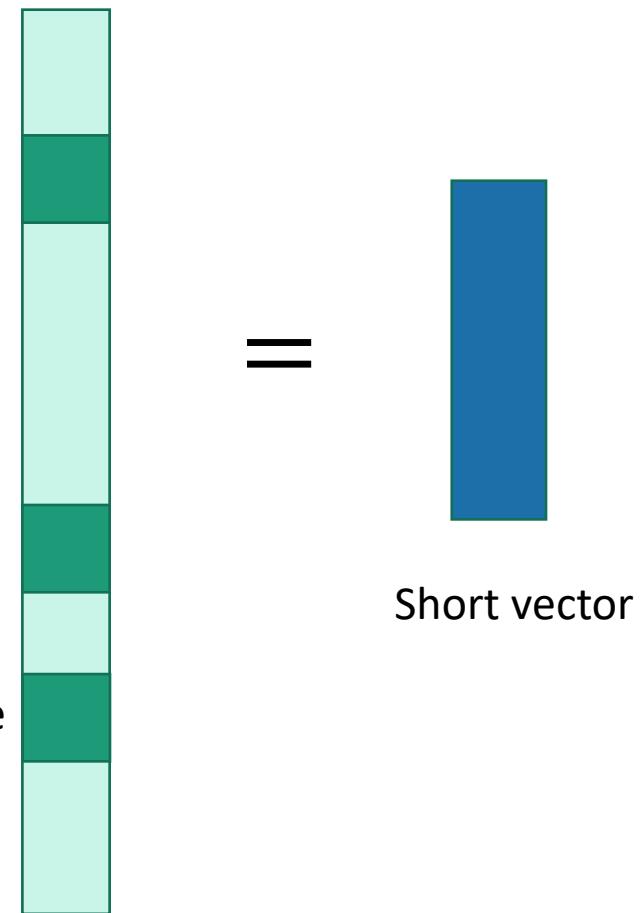


Goal: Given the **short vector**, recover the **long sparse vector**.

But why should this be possible?

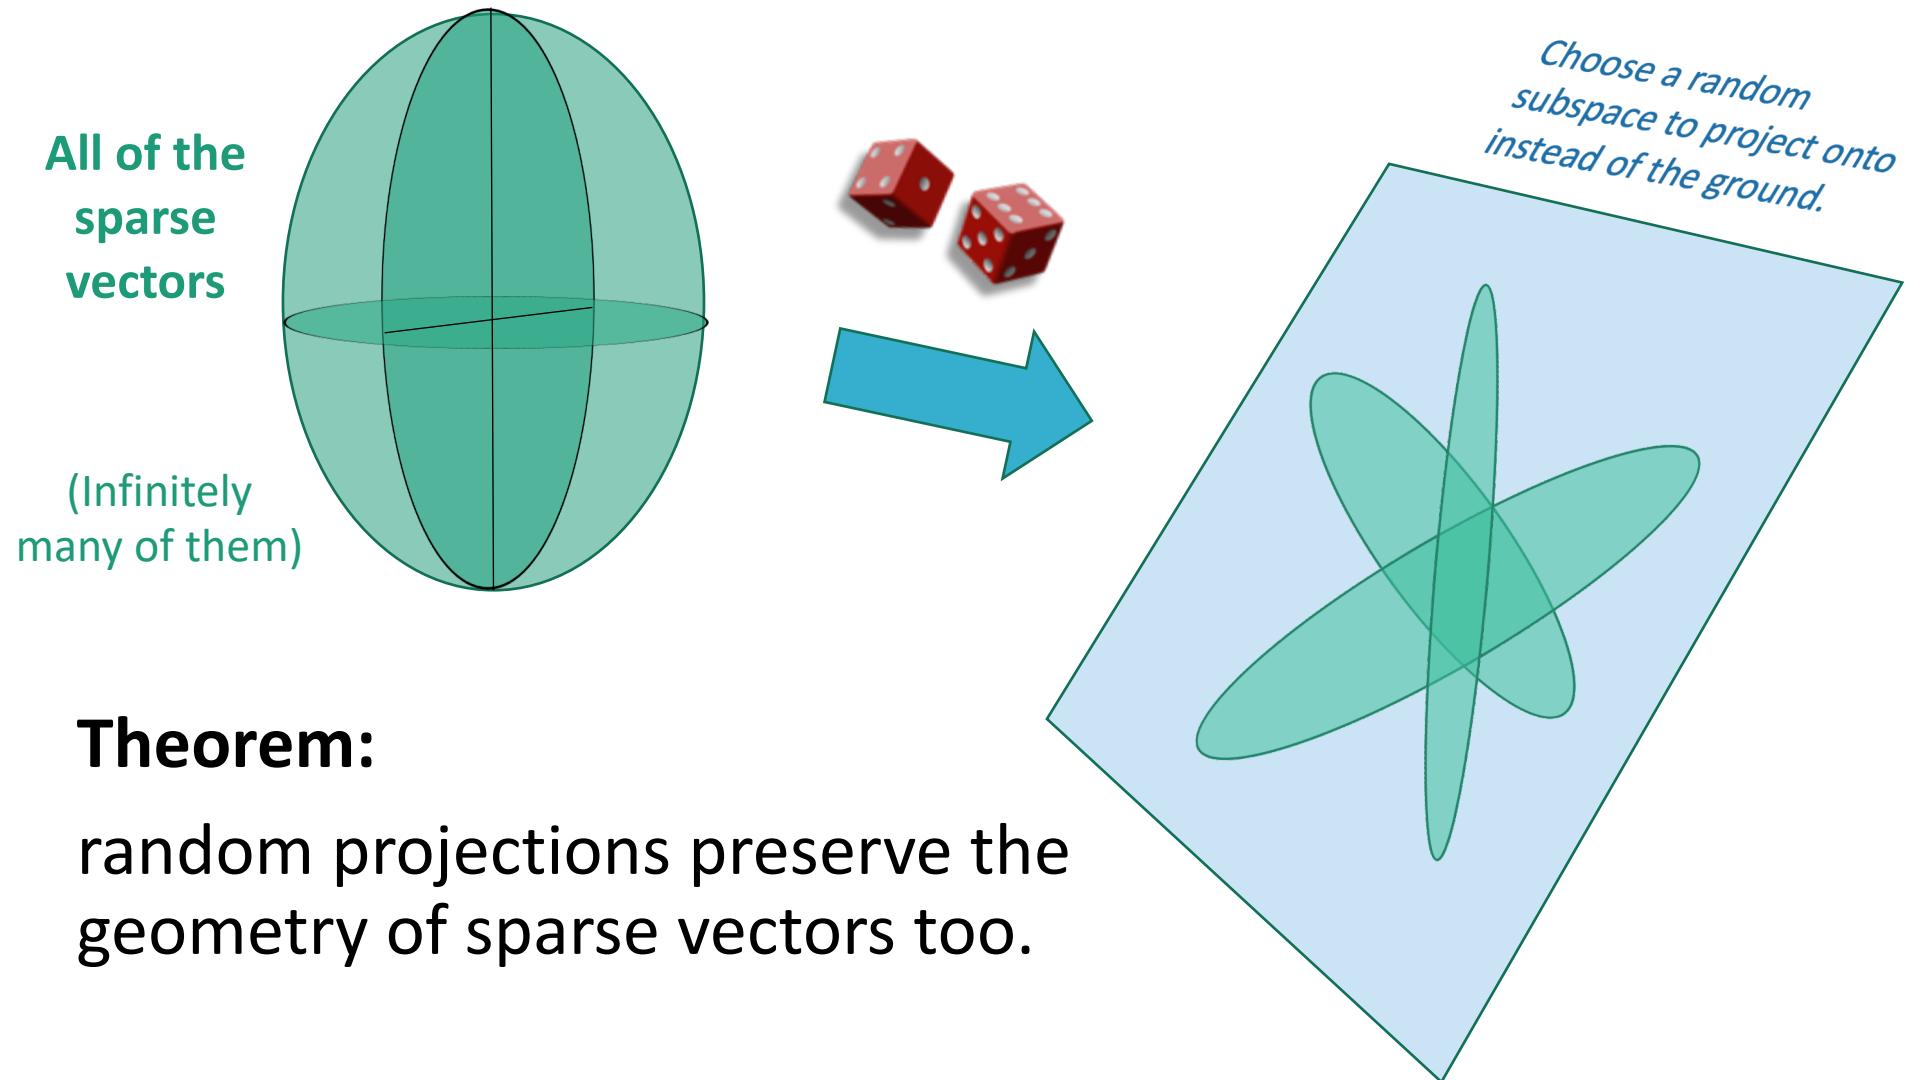
- There are tons of long vectors that map to the short vector!

Random short fat matrix



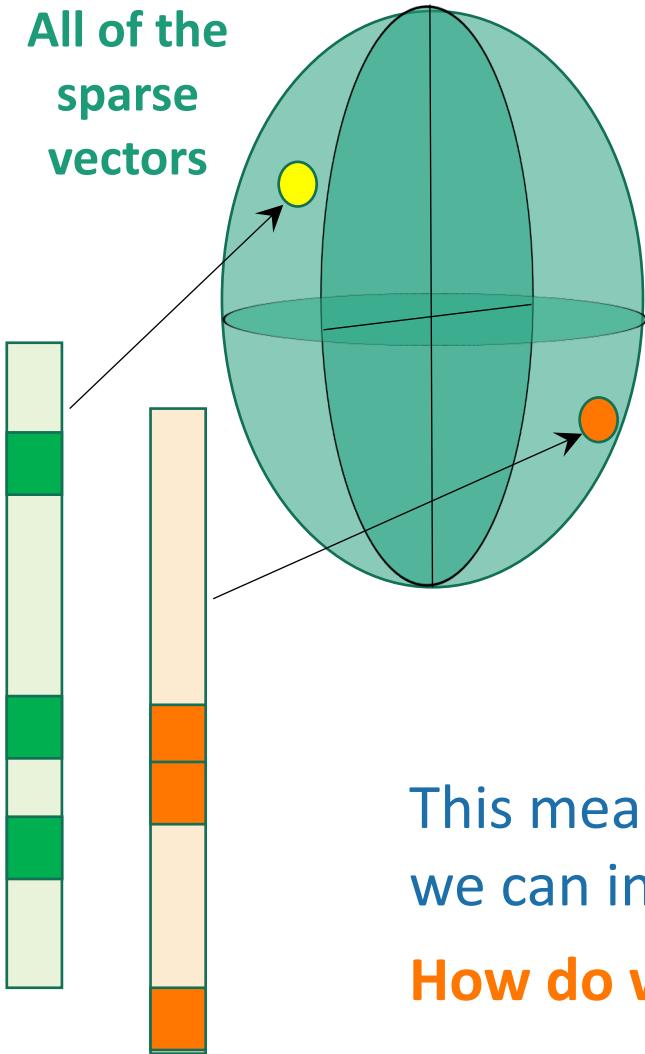
Goal: Given the **short vector**, recover the **long sparse vector**.

Back to the geometry



If we don't care about algorithms,
that's more than enough.

All of the
sparse
vectors

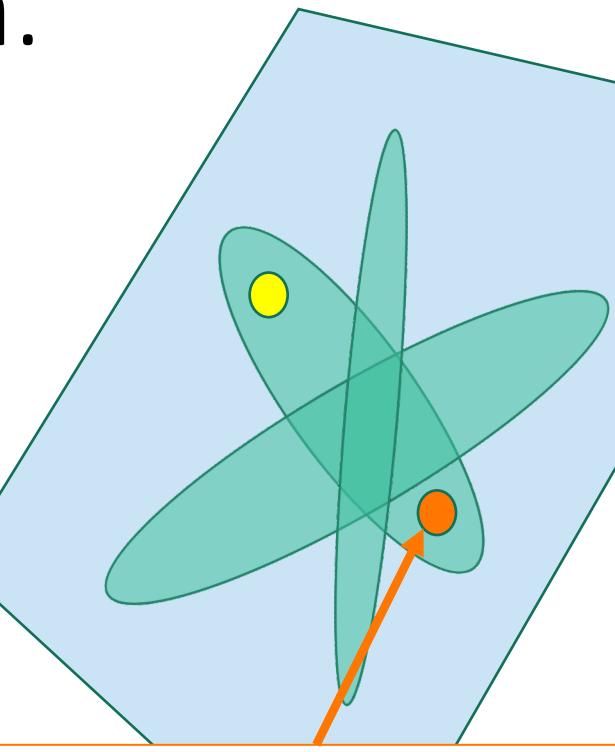


Multiply by

Random short
fat matrix

This means that, in theory,
we can invert that arrow.

How do we do this efficiently??



There may be tons of vectors
that map to this point, but only
one of them is sparse!

Goal: Given the **short vector**,
recover the **long sparse vector**.

An efficient algorithm?

What we'd like to do is:

Minimize number of
nonzero entries in x

s.t.

$$Ax = y$$

This norm is the sum
of the absolute values
of the entries of x

This isn't a
nice function

Instead:

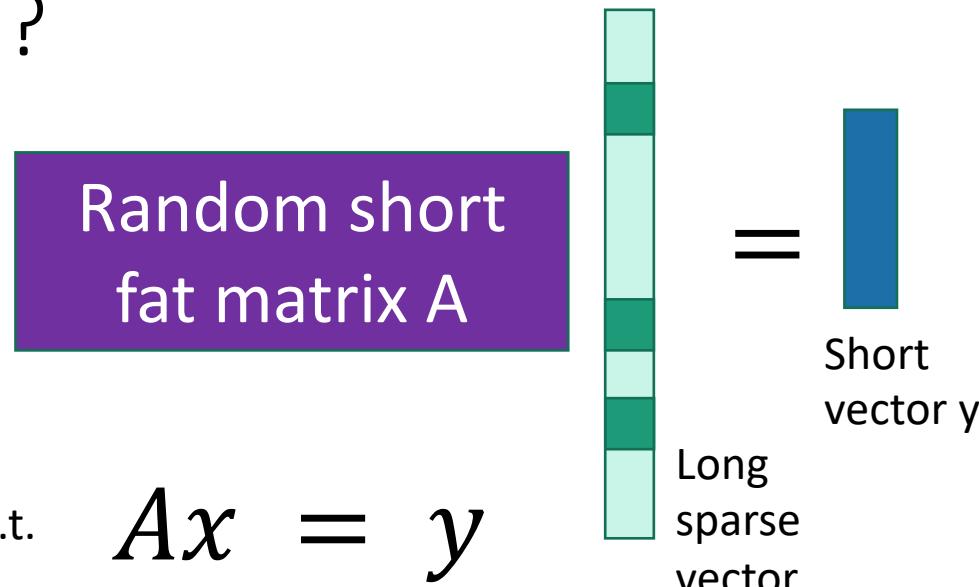
Minimize $\|x\|_1$

Problem: I don't know
how to do that efficiently!

s.t.

$$Ax = y$$

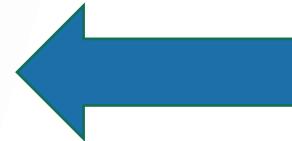
- It turns out that because the geometry of sparse vectors is preserved, this optimization problem **gives the same answer**.
- We can use **linear programming** to solve this quickly!



Today

A few gems

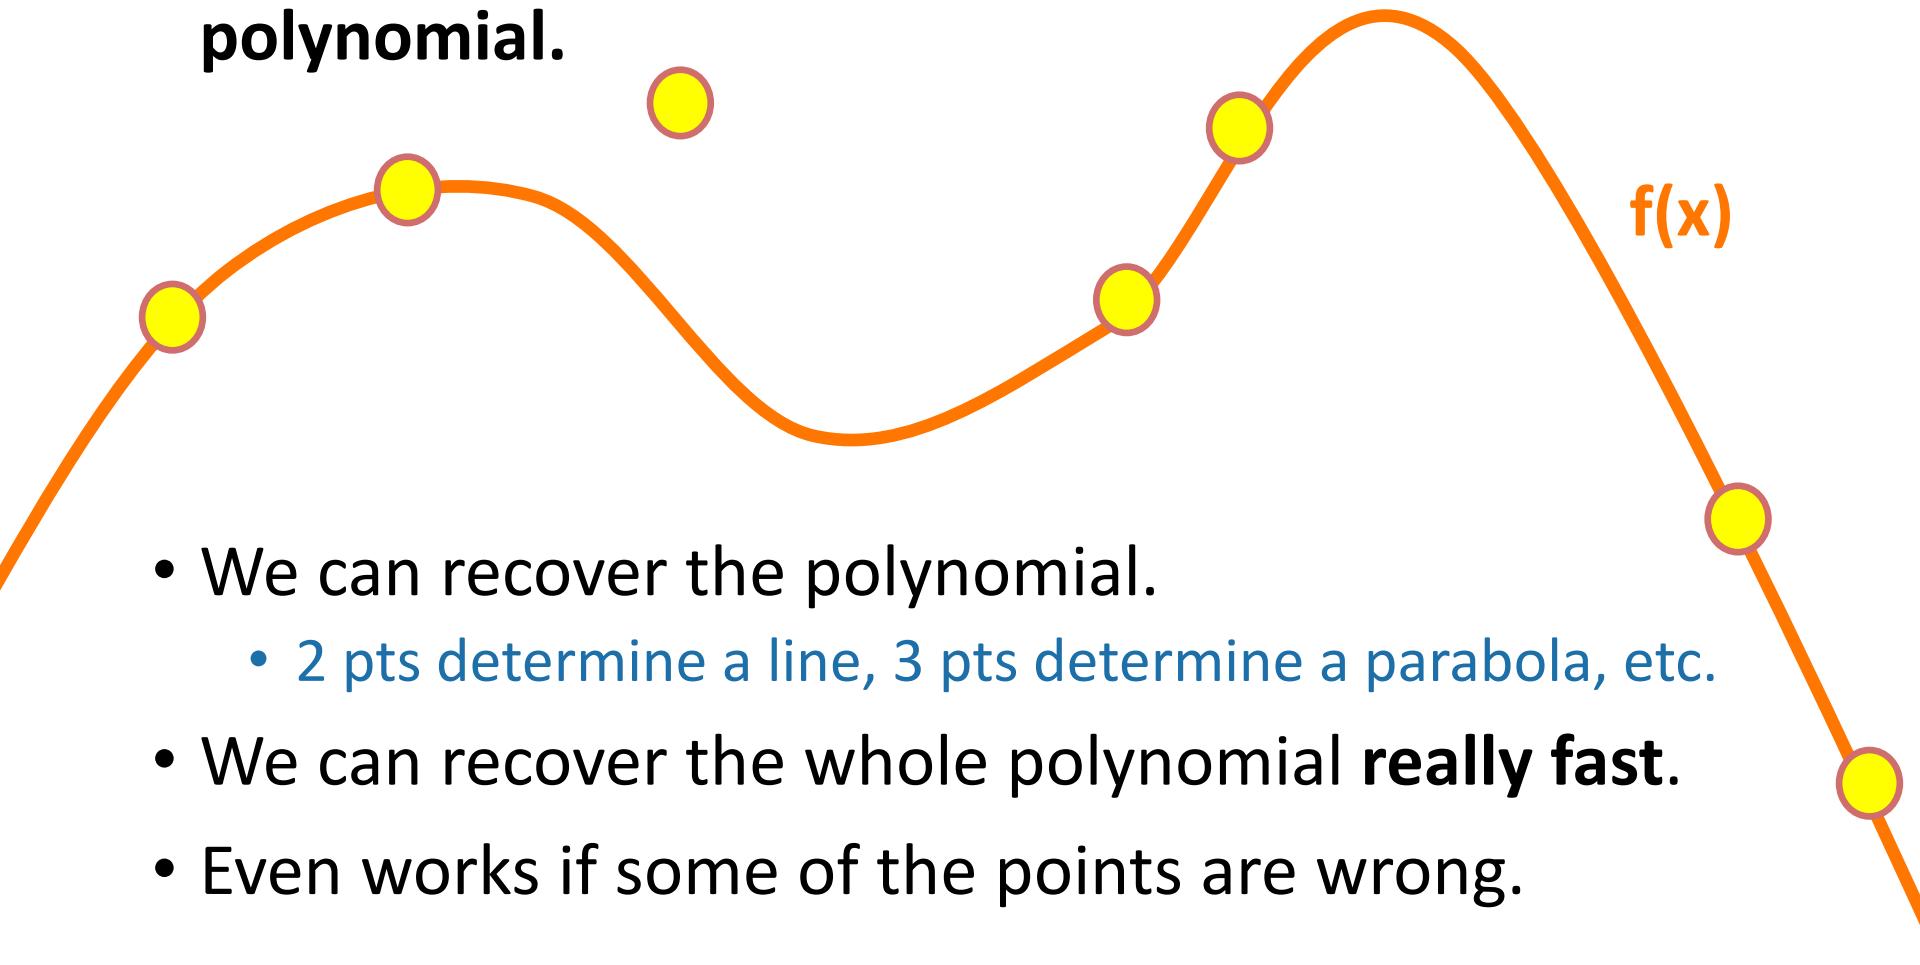
- Linear programming
- Random projections
- Low-degree polynomials



Another very useful trick

Polynomial interpolation

- Say we have a few evaluation points of a **low-degree polynomial**.



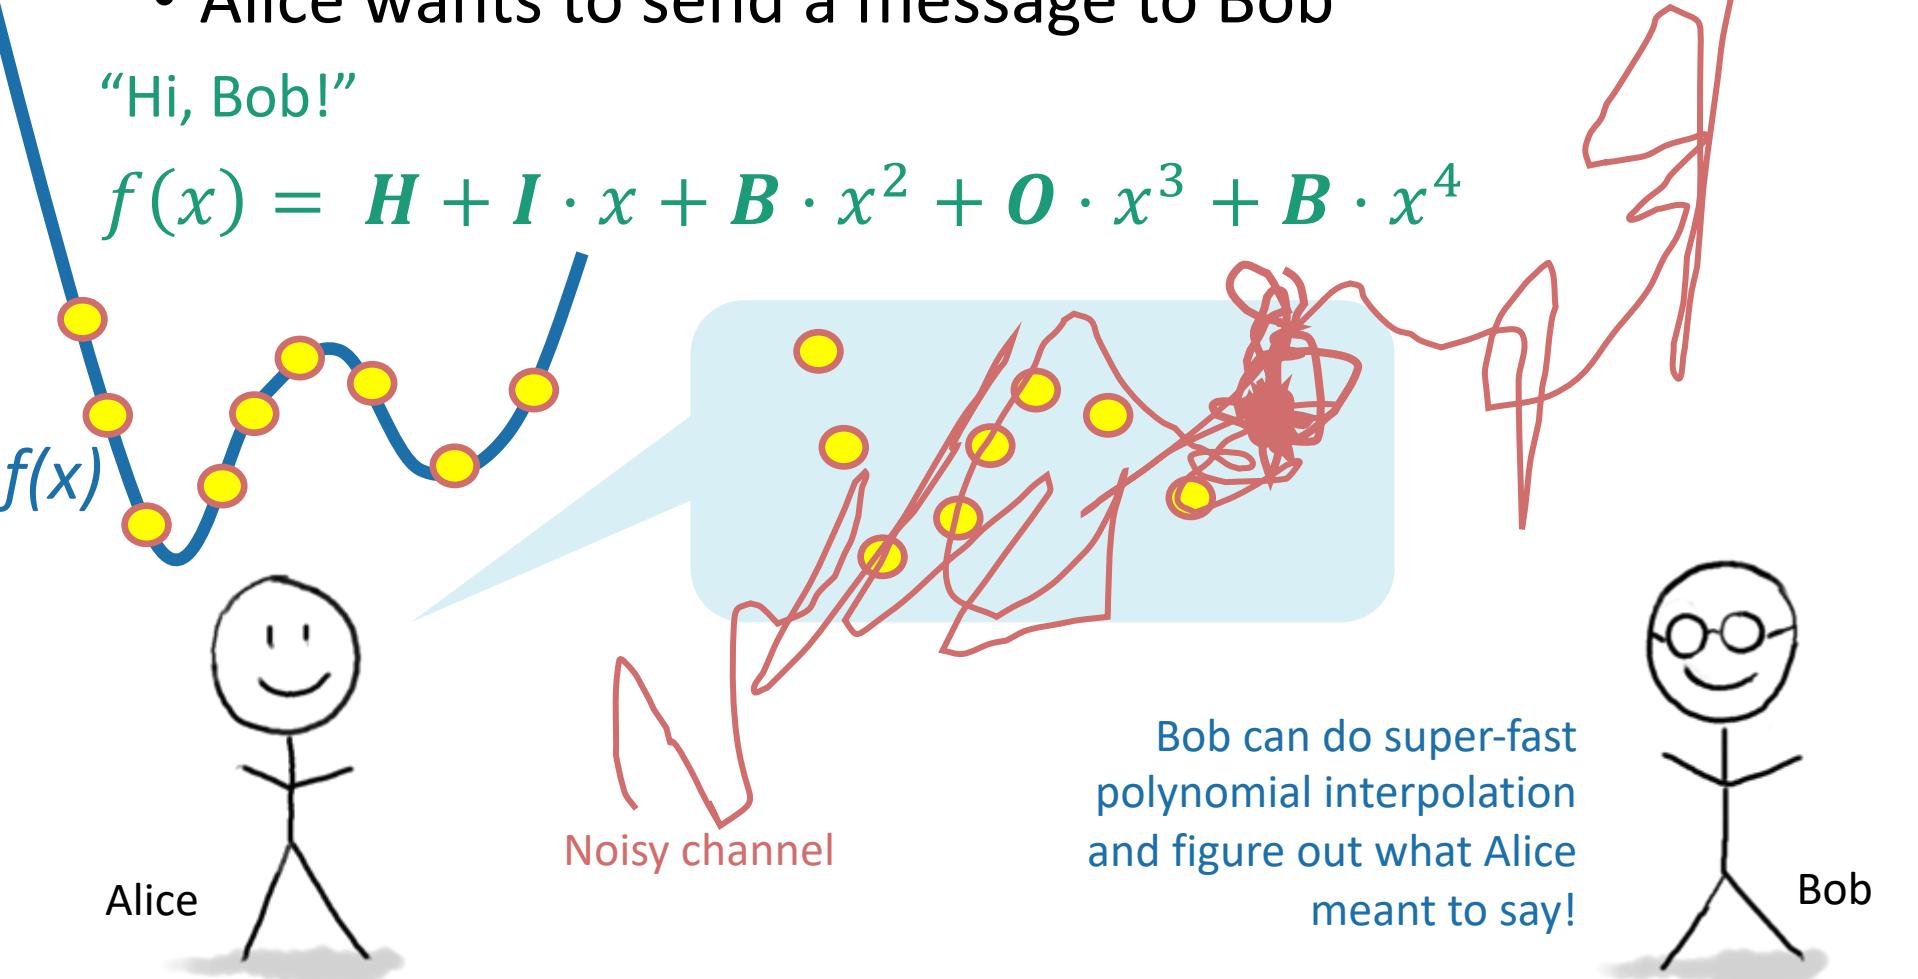
- We can recover the polynomial.
 - 2 pts determine a line, 3 pts determine a parabola, etc.
- We can recover the whole polynomial **really fast**.
- Even works if some of the points are wrong.

One application: Communication and Storage

- Alice wants to send a message to Bob

"Hi, Bob!"

$$f(x) = H + I \cdot x + B \cdot x^2 + O \cdot x^3 + B \cdot x^4$$



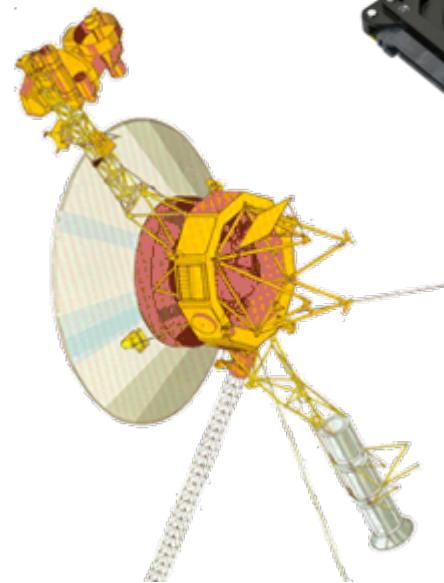
Bob can do super-fast
polynomial interpolation
and figure out what Alice
meant to say!

Alice

Bob

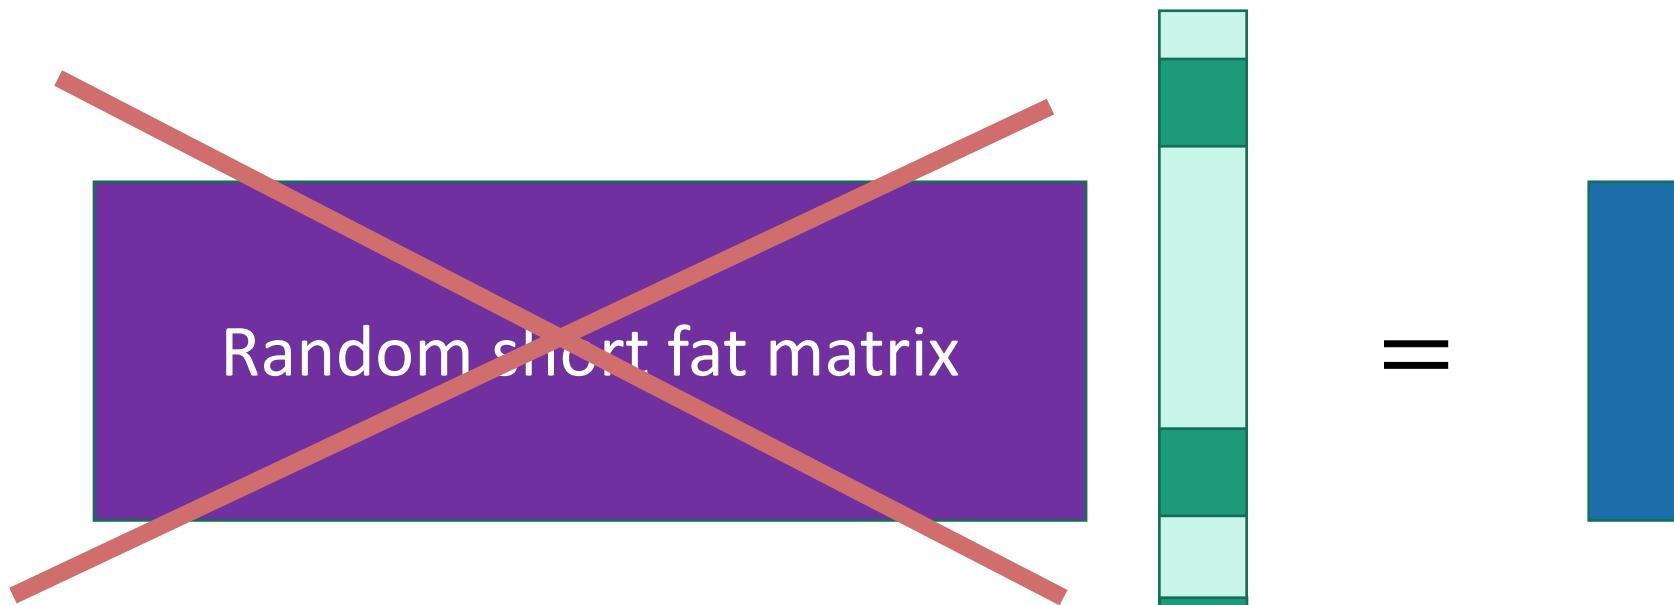
This is used in practice

- It's called “Reed-Solomon Encoding”



Another application:

Designing “random” projections that are better than random



The matrix that treats the big long vector as Alice's message polynomial and evaluates it REALLY FAST at random points.

- This is still “random enough” to make the LP solution work.
- It is much more efficient to manipulate and store!

Today

A few gems

- Linear programming
- Random projections
- Low-degree polynomials



To learn more:

CS168, CS261, ...

CS168, CS261,
CS265, ...

CS168, CS250, ...

What have we learned?

CS161



Tons more cool
algorithms stuff!

To see more...

- Take more classes!
- Come hang out with the theory group!
 - In person, once we can hang out in person again ...
 - Theory lunch, most Thursdays at noon (remote for now).
 - Join the theory-seminar mailing list for updates.



theory.stanford.edu

Stanford theory group (circa 2017):
We are very friendly.

A few final messages...

Thanks to our course coordinator
Amelie Byun!

- Amelie has been making all the logistics work behind the scenes.



Thanks to our superstar CAs!!!

tell them you appreciate them!



Albert Zuo



Amy Kanne



Avery Wang



Caci Jiang



Carrie Wu



Changyu Bi



Reyna Hulett



Chenru Liu



Geng Zhao



Ian Tullis



Jerry Qu



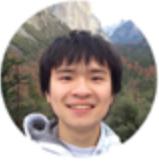
Jesus Cervantes



Jiazheng Zhao



John Sun



Nick Lai



Ofir Geri



Qile (Suyie) Zhi



Rose Li



Teresa Noyola



Trey Connelly



Weiyun (Anna) Ma



Wilhem Kautz

4.

THANKS
to you!!!!

